

# Introduction

Daniel Huynh



- Graduate from Polytechnique (X2016) and X-HEC Data Science for Business
- Former intern at Microsoft France in the office of the CTO
- CEO of Mithril Security, a software solution to secure company external data sharing

MACHINE LEARNING

CRYPTOGRAPHY

PRIVACY PRESERVING ML

# Our agenda for this Common Track

Context and cloud data threads

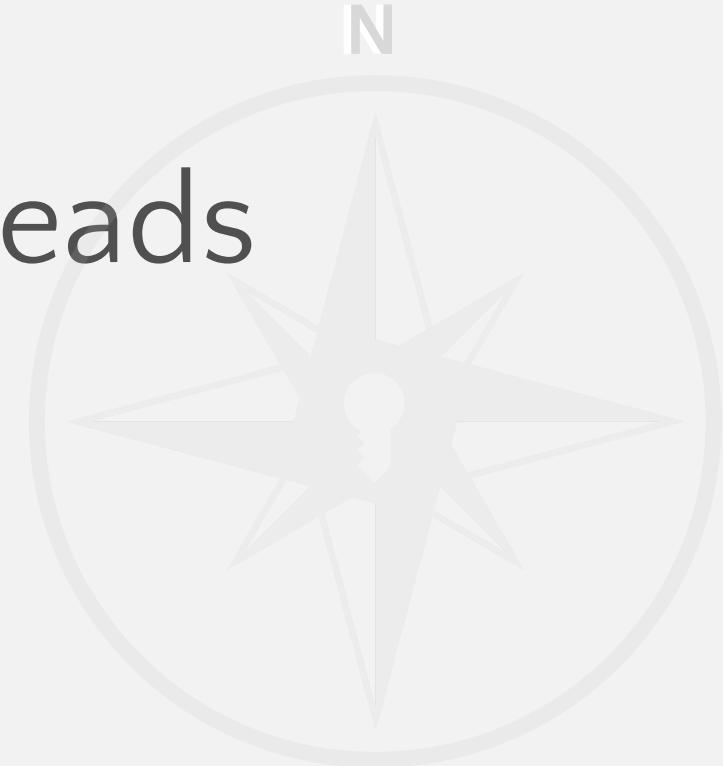
Core mitigation techniques

- Confidential Computing (CC)
- Homomorphic Encryption (HE)
- Secure Multi-Party Computing (MPC)

Other techniques to consider (if time permits)

- Differential Privacy (DP)

# Context and cloud data threads



# Data breaches need to be taken seriously

LifeLabs cyberattack one of 'several wake-up calls' for e-health security and privacy

December 19, 2019



"What's interesting about the large push for electronic patient health-care information that you put online is that a lot of these [organizations are not designed to withstand attacks](#)."

Canada's Desjardins reports costs of \$53M related to data breach

August 13, 2019



["Unauthorized use of internal data by an employee led to breach of personal information including social insurance number, address and details of banking habits."](#)

Anthem is warning consumers about its huge data breach. Here's a translation.

March 6, 2015



"The indications are that they gained access to Anthem's data by [stealing the network credentials of at least five employees with high-level IT access](#)."

# Data breaches are costly

**7.27% decrease**

Is the average price of a company disclosing a data breach<sup>1</sup>

**\$3.92M avg.**

The Cost of a data breach to an enterprise. (Up 12% in 5 years).<sup>2</sup>

**\$6.5M avg.**

The healthcare cost of a breach (over 60% more than other industries in the study).<sup>2</sup>

**\$42M**

Projected in losses for 1 million records losses to cost<sup>2</sup>

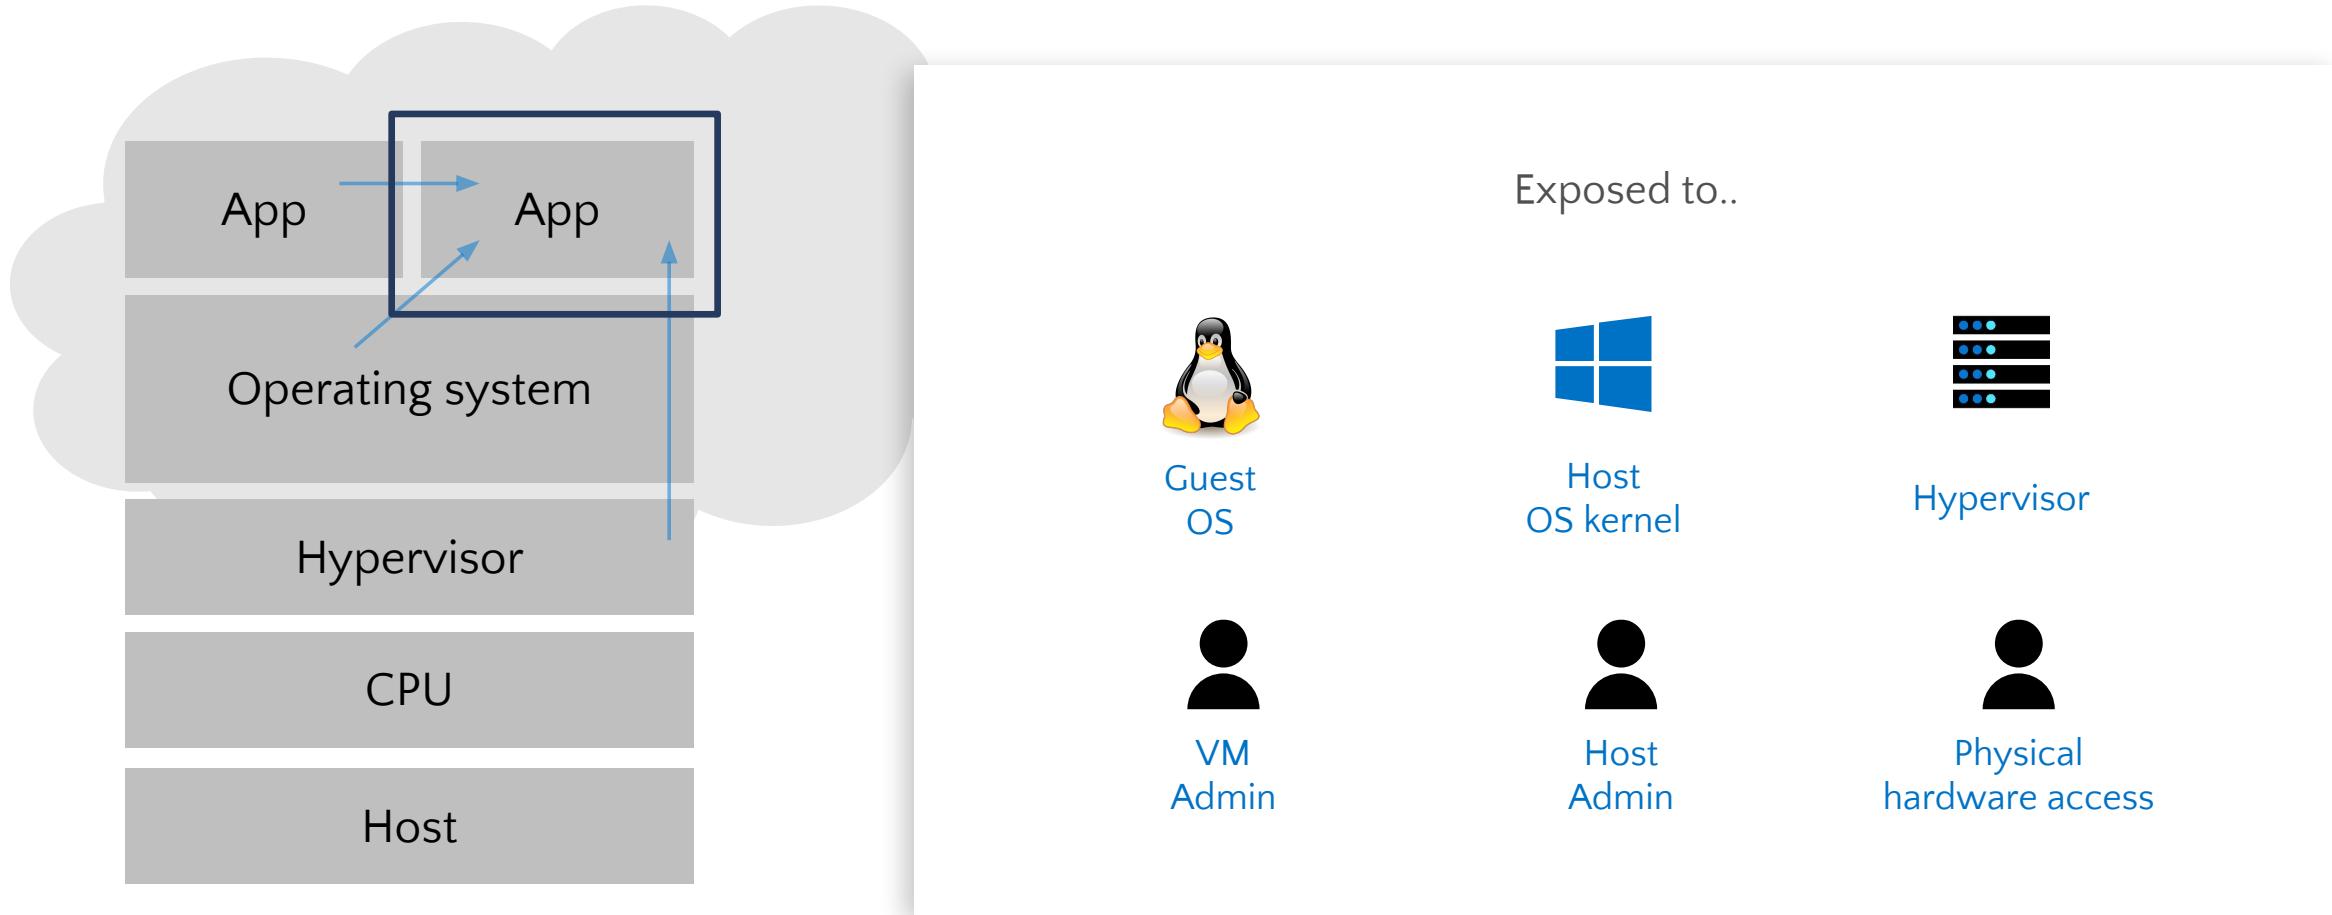
**\$388M**

Projected losses to cost >50 million records<sup>2</sup>

Sources: 1. [Comparitech](#); 2. [IBM](#)



# Why do we need to also protect “data in use”?



# Confidential Computing (CC)

For your security, regulatory, and compliance needs

The ability to store, transport, and act on compute workloads without compromising privacy of data and intellectual property



# How to control the (most sensitive) data throughout its complete lifecycle?

Confidential Computing technologies are helping us evolve from computing in the clear to protecting data while in use, reducing the need for trust in HW/SW stacks & operators

Existing Encryption	New
	
<p><b>Data at rest</b></p> <p>Encrypt inactive data when stored in blob storage, database, etc.</p> <p>The industry <u>moved from disks in the clear to encrypted disks, with managed keys</u></p>	<p><b>Data in transit</b></p> <p>Encrypt data that is flowing between untrusted public or private networks</p> <p>Evolved from <u>browsing/moving data in the clear</u> (HTTP), to <u>encrypting data</u> (HTTPS / TLS)</p>
	<p><b>In use</b></p> <p>Protect/Encrypt data that is in use, while in RAM and during computation</p> <p>Evolving from computing in the clear to Trusted Execution Environments (TEEs), like Intel SGX</p>

# Who do you need to trust?

App software vendor

Trusted today on-premises

Use open source or build own software.

Infrastructure provider

Trust a cloud provider, or

Manage own on-premises datacenters.

Hardware vendor

Trusted today on-premises

Build own hardware.

## Attack surface

Application runtime &  
dependencies, many open  
source packages

Other tenants  
running on the host,  
infrastructure services

Observable side  
channels in hardware

## Software TCB



Application



VM OS



Hypervisor



Host OS



Hardware

## Trusted actors



App Admin



VM Admin



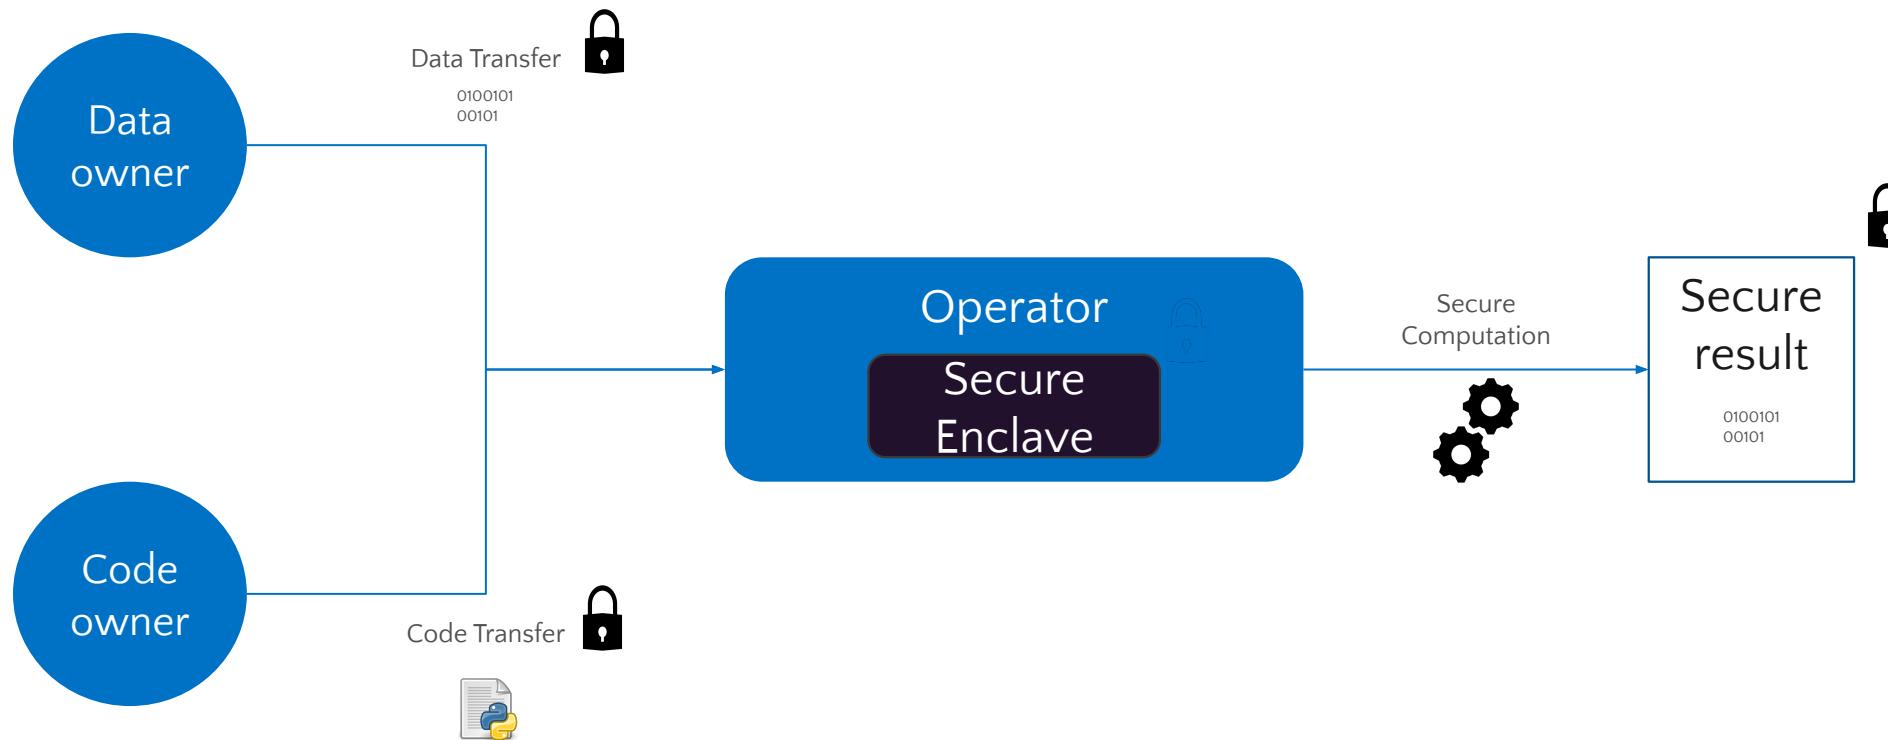
Host Admin



Physical  
hardware access

# High level view of Confidential Computing with Intel SGX

**Confidential computing allows execution of code we do not see on data we do not see**



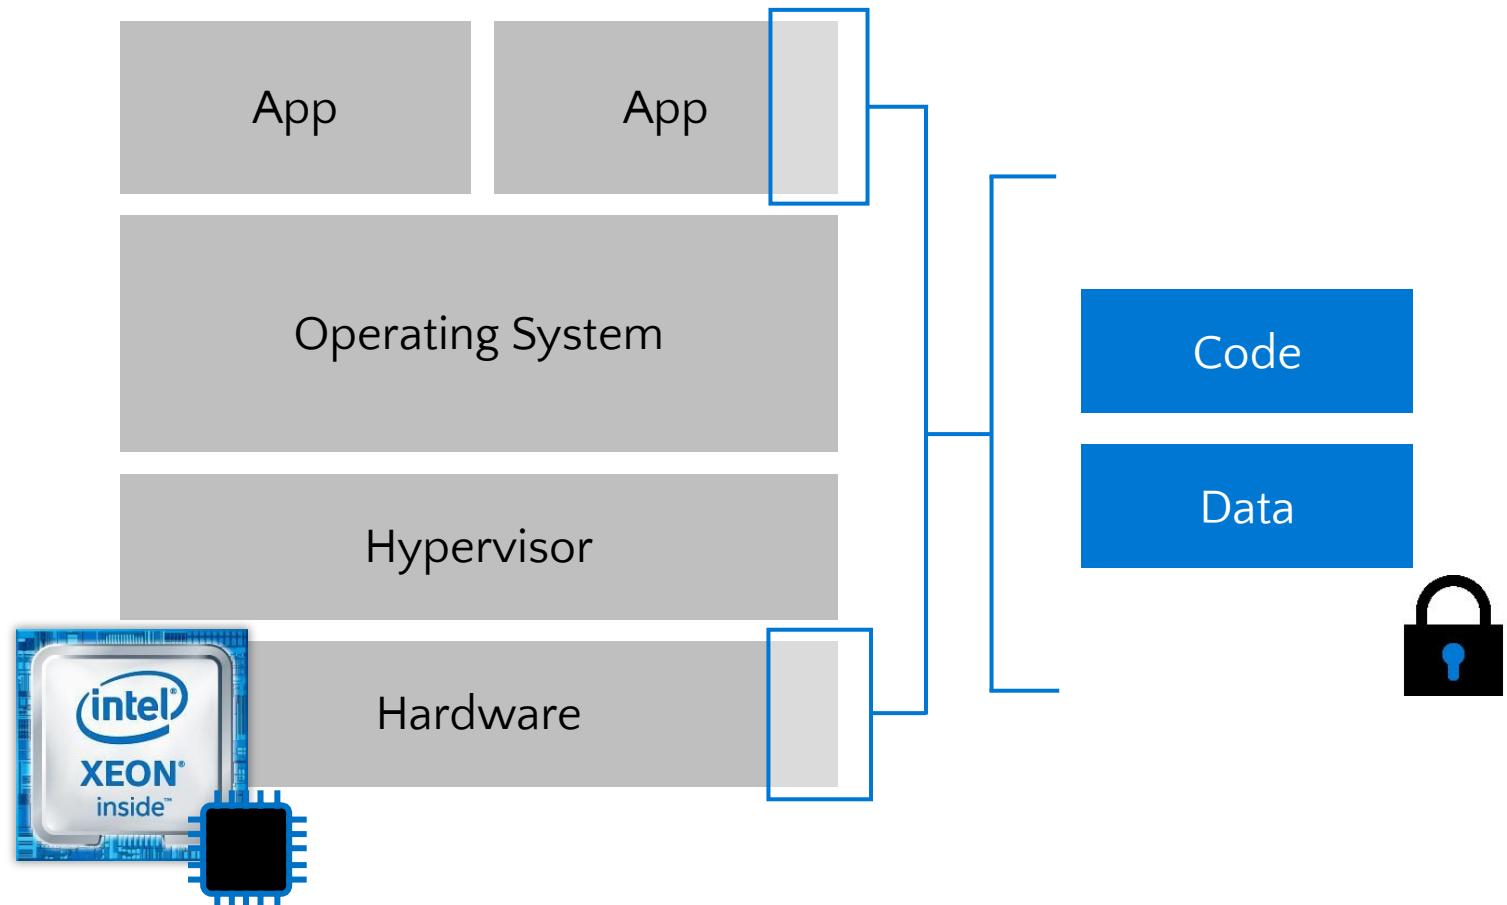
# How Intel SGX hardware protects data in use

**Intel SGX Goal:** Minimize attack surface to CPU

New hardware architecture

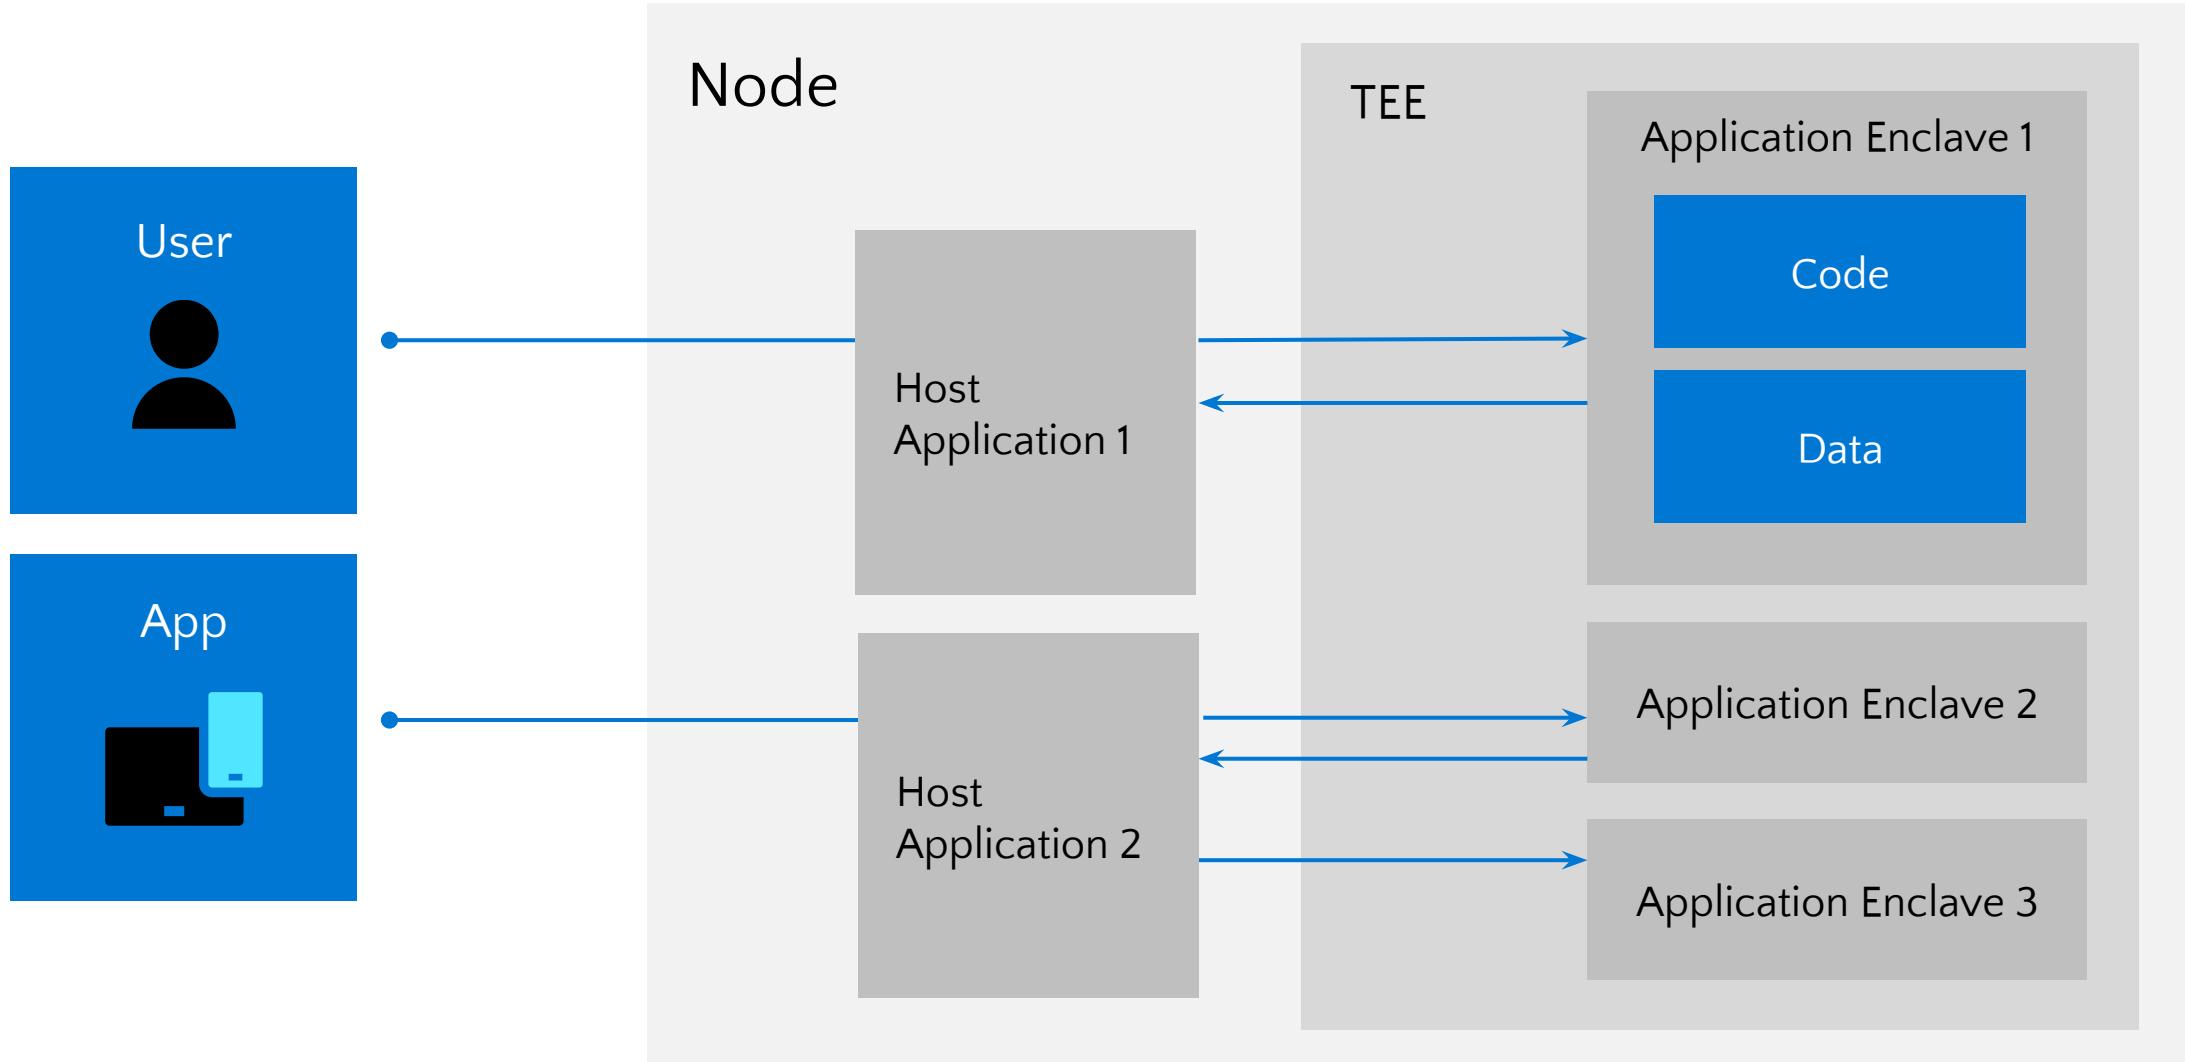
New instructions to set aside private regions (“protected containers”) of code and data

Data is only ever in the clear within the protected memory and encrypted if pushed off of CPU



# An enclave-based application model

## Development experience and software

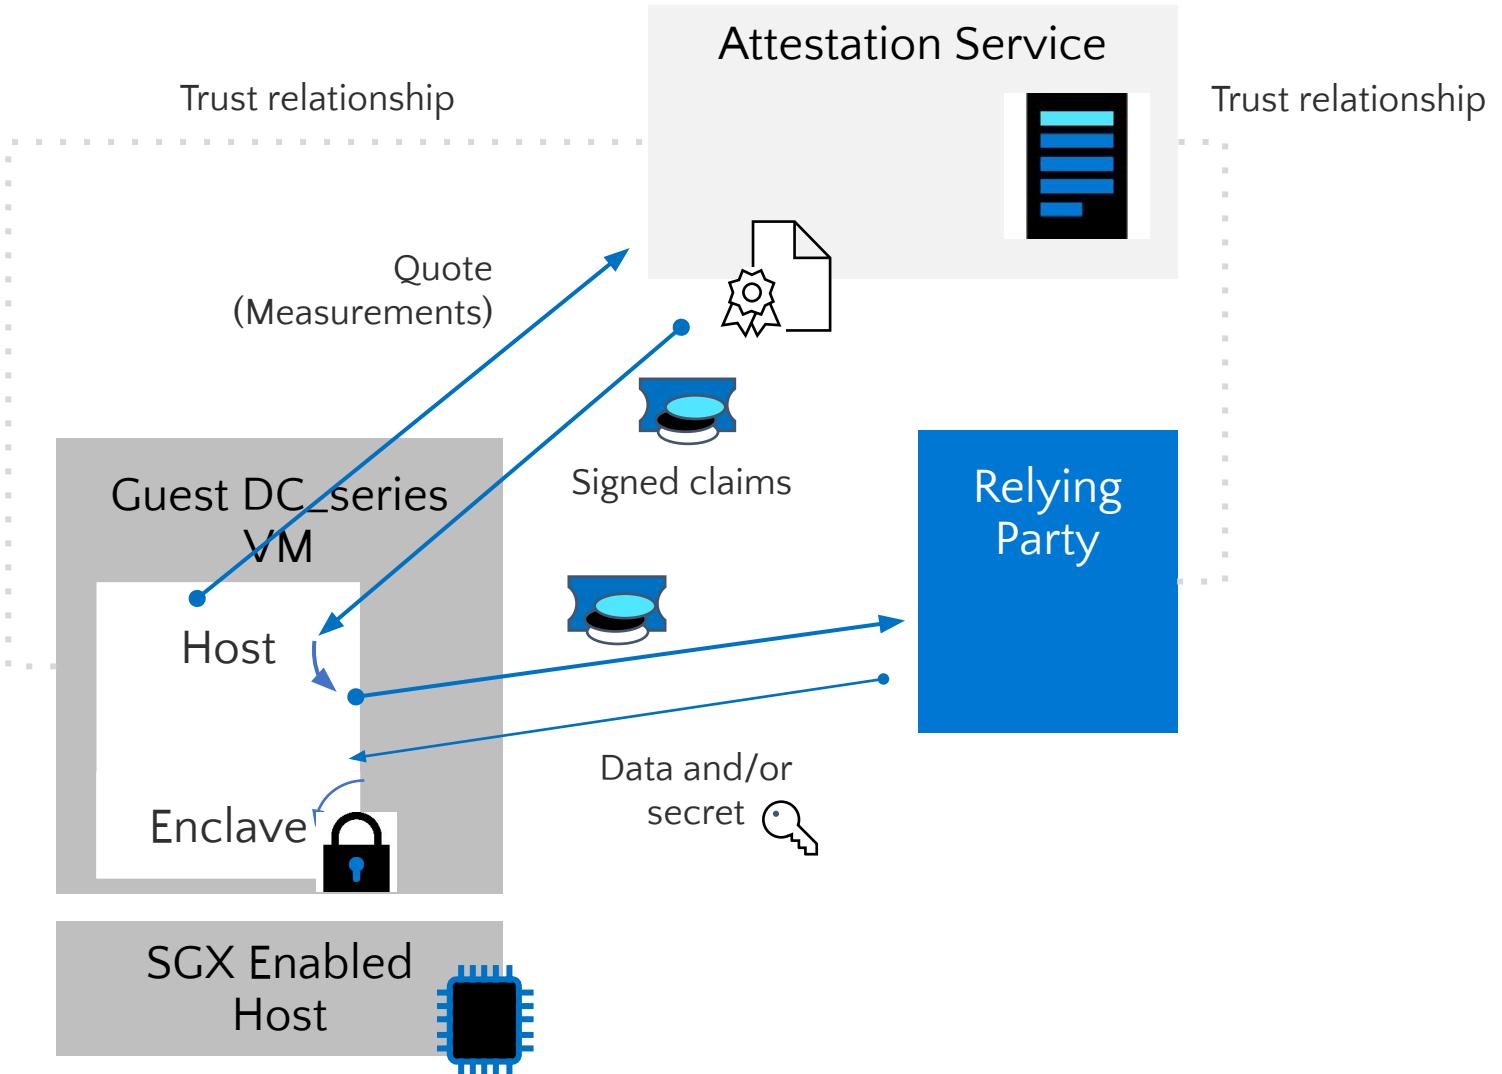


# Azure Attestation Service

## Azure Confidential Computing

### Service for Confidential Computing

1. Quote provide proofs:
  - Code runs in genuine SGX enclave
  - Enclave version and owner is as expected
  - Arbitrary enclave supplied data is as expected
2. Attestation service attests to hardware, rooted in Intel chain of trust, and issues signed claims as a security token
3. Relying party, e.g. another enclave, a cloud service, etc. is presented with security token with certs chained to CPU
4. Security token is used to release data and/or secrets to the application enclave



# Confidential computing can help

## Use cases

1

Make your [applications more secure](#) by leveraging confidential computing VMs and the Open Enclave SDK

2

Enable [Customer workloads](#), confidential blockchain, secrets storage and processing, analytics, ML training and inferencing, datastores, IoT

3

Preserve privacy and sensitive customer data across organizations with [multi-party dataset analytics and machine learning](#), both training and inference.

# Three abstraction layers for enclaves

## Development experience and software

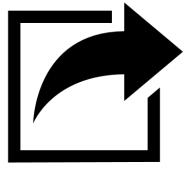


### New/Refactor

Understand fundamentals, create a trusted portion of your application to run inside the enclave/TEE with “the building blocks”

Open Enclave SDK  
(OE SDK)

Confidential Consortium Framework (CCF)



### “Lift and shift” model

Application just works inside secure hardware with no/minimal changes

Secure containers/LibOS Strategy

Open source, such as SGX-LKL

Leverage one of the ISVs available in the Marketplace (e.g. Fortanix, Anjuna)



### Confidential workloads using PaaS services

It's taken care of by services already used

SQL Azure DB Always Encrypted

Confidential ML

Confidential AKS

Confidential Blockchain

# Open Enclave SDK

## Development experience and software: New/Refactor

Driving towards a consistent API surface around  
enclaving abstraction, supporting portability across  
enclave types and flexibility in architecture

Open sourced by Microsoft in October 2018 under the  
MIT license

<https://github.com/openenclave/openenclave>

More than 6,000 commits and 16 versions since the original release date

More than 70 contributors, 150 forks, 400 stars and 1,200 pull-requests

Very active project with 25-100 commits per week

### Core functionalities

- Enclave creation and management
- Expressions of enclave measurement and identity
- Mechanisms for defining call-ins and call-outs
- and the data marshalling associated with them
- System primitives exposed by enclave runtime,
  - such as thread and memory management
  - Sealing functions
  - Attestation functions
- Runtime and cryptographic library support

The screenshot shows the GitHub repository page for 'openenclave / openenclave'. Key statistics at the top include 5,863 commits, 86 branches, 0 packages, 15 releases, and 71 contributors. The repository is licensed under MIT. Below the stats, a list of recent commits is shown, each with a small icon, the author, a brief description, and the time since the commit. The commits are from various contributors like jenkins, 3rdparty, cmake, common, debugger, devex, docs, enclave, host, include, libc, libcxx, pkgconfig, prereqs, samples, scripts, syscall, tests, tools, .clang-format, .cspellignore, and .gitattributes. The commits range from 10 hours ago to 2 years ago.

### Extending support to:

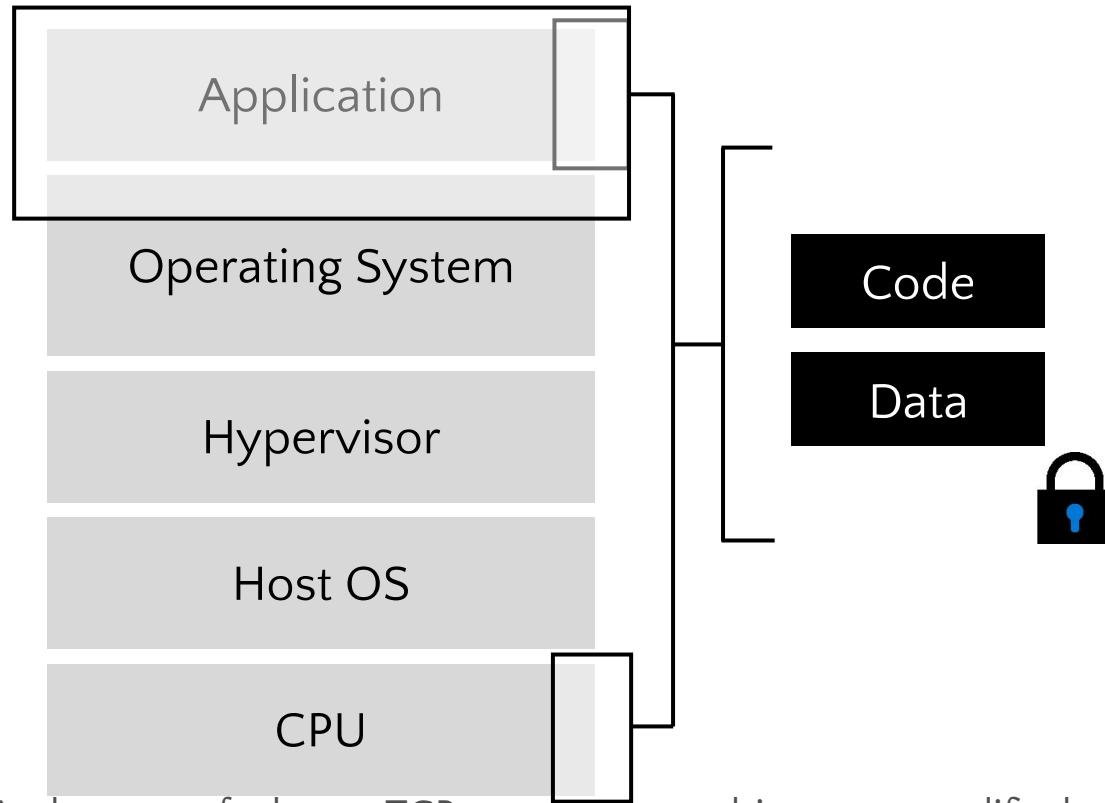
- Operating System: Linux, Windows, OP-TEE OS\*
- Trusted Execution Environments: Intel SGX, Arm TrustZone
- Runtime Libraries: C/C++
- Cryptographic Libraries: mBedTLS

\*OP-TEE OS: <https://www.op-tee.org/>

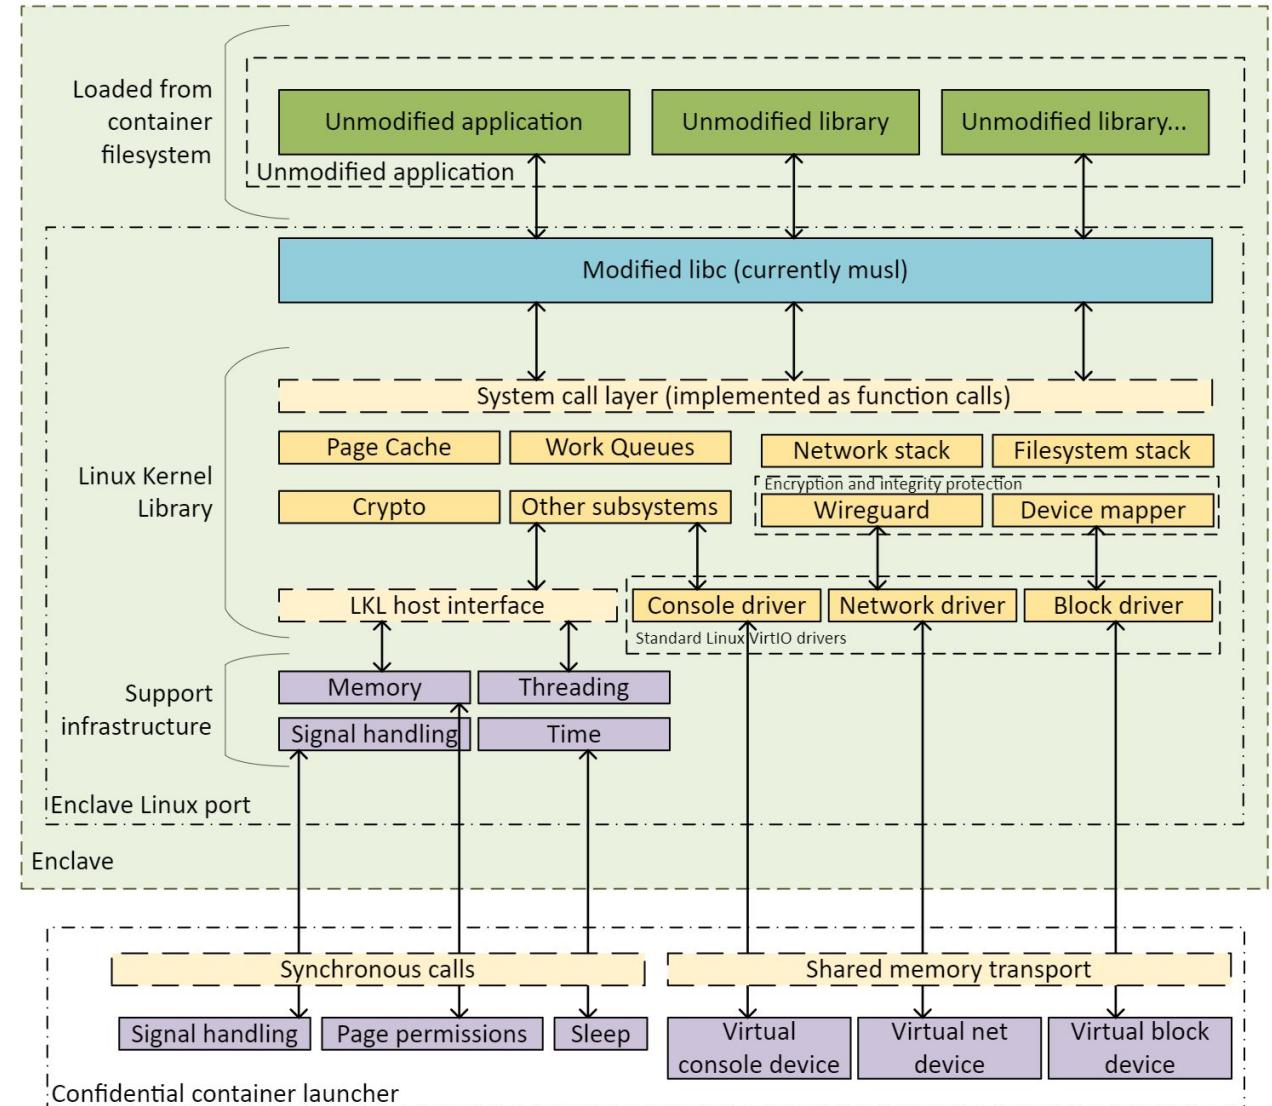
# Make your code confidential without changing a line!

Development experience and software: “Lift and shift” model

## Confidential Linux containers with SGX-LKL

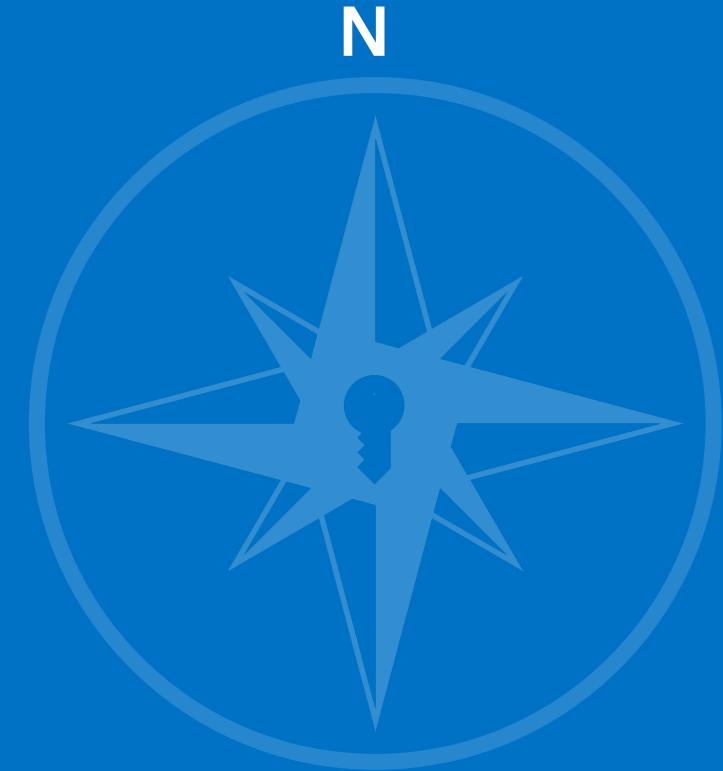


At the cost of a larger TCB, we can run arbitrary unmodified applications, such as TensorFlow



# Demo

Image classification with Graphene SGX



# Use Confidential Computing with Kubernetes

## Development experience and software

An additional layer of data protection for the Kubernetes workloads with the code running on the CPU with secure hardware enclaves

Using the Open Enclave SDK for confidential computing in code with ACC DC\_series VMs

More info: [Bringing confidential computing to Kubernetes](#)

The oe-sgx device plugin (alpha) surfaces the usage of Intel SGX's Encrypted Page Cache (EPC) RAM as a schedulable resource for Kubernetes

Allow to schedule pods and containers to use the Open Enclave SDK, and to have access to a TEE by defining a limit on the specific EPC memory that is advertised to the Kubernetes scheduler by the device plugin

See [Using SGX with Kubernetes \(AKS\\_Engine\)](#)

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: oe-deployment
5  spec:
6    selector:
7      matchLabels:
8        app: oe-app
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         app: oe-app
14     spec:
15       containers:
16         - name: <image_name>
17           image: <image_reference>
18           command: <exec>
19         resources:
20           limits:
21             openenclave.io/sgx_epc_MiB: 64
```

# Multi-party machine learning use case



## Goals

Data and Model/IP confidentiality from other parties and/or cloud provider

Control integrity of ML code running against data set from malicious parties



## Solution

Run ML modelling process (data input, pre-processing, model training, etc.) in a trusted execution environments

Decrypt data only within a TEE to protect confidentiality at centralized point

Run code that is attestable at centralized point

## Properties

Confidentiality and integrity of data & models (from other parties)  
Transparency (only certain ML code can access the data)  
Auditability (logs of activities)

## Protocols

Agree on machine learning algorithm/code to use  
Upload code  
Verify enclave setup (hardware and software identity)  
Send data and keys to system  
Receive computed output

Multi-party machine learning:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/paper.pdf>

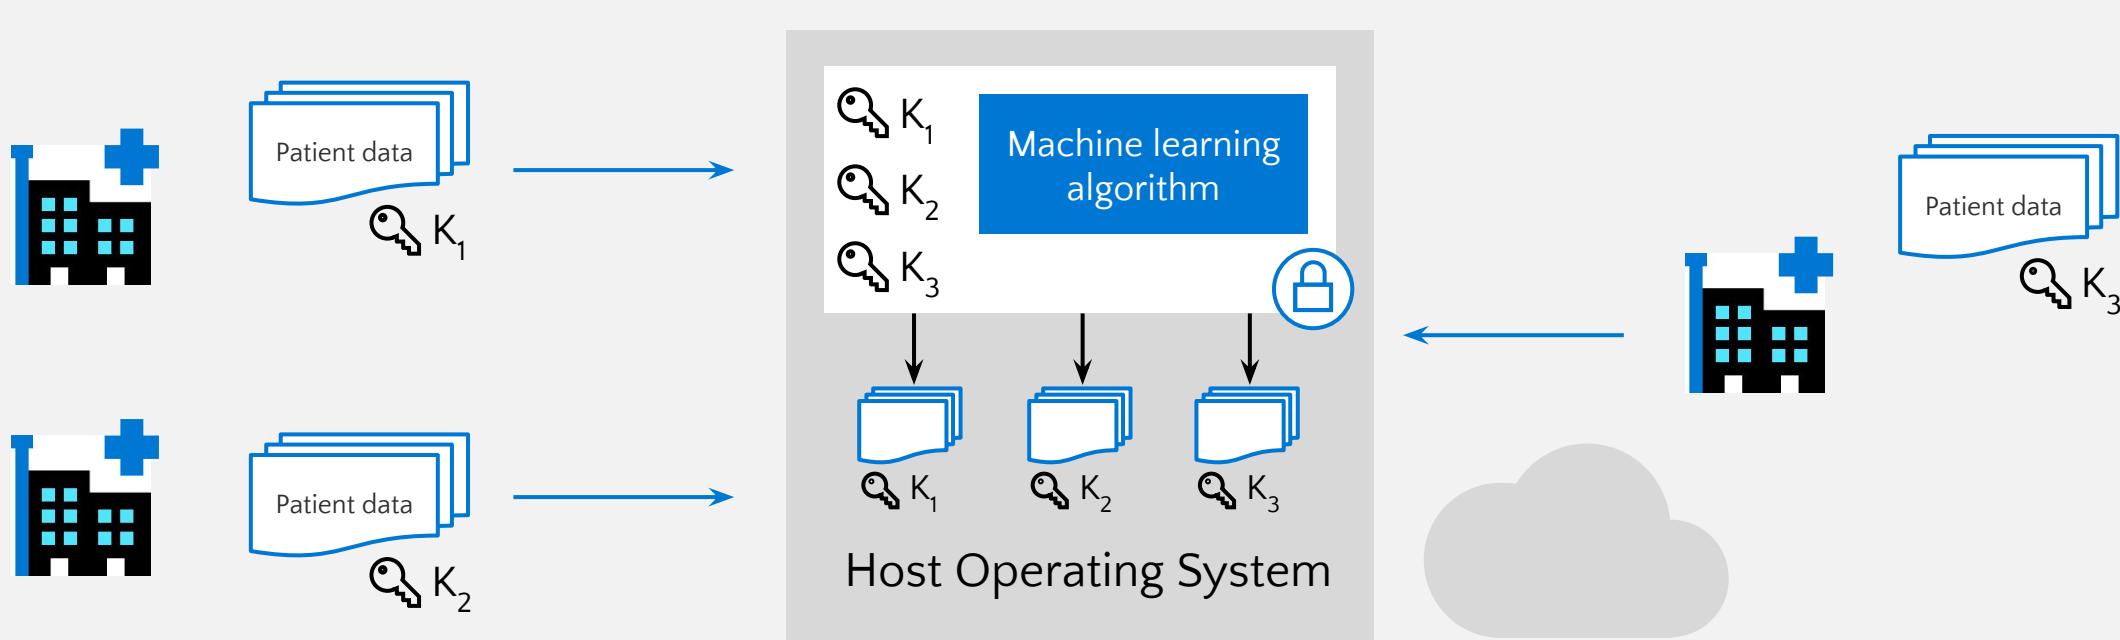
# Healthcare scenarios

Development experience and software: Confidential workloads (using PaaS services)

Partnered health facilities contribute private patient health data sets to train a ML model

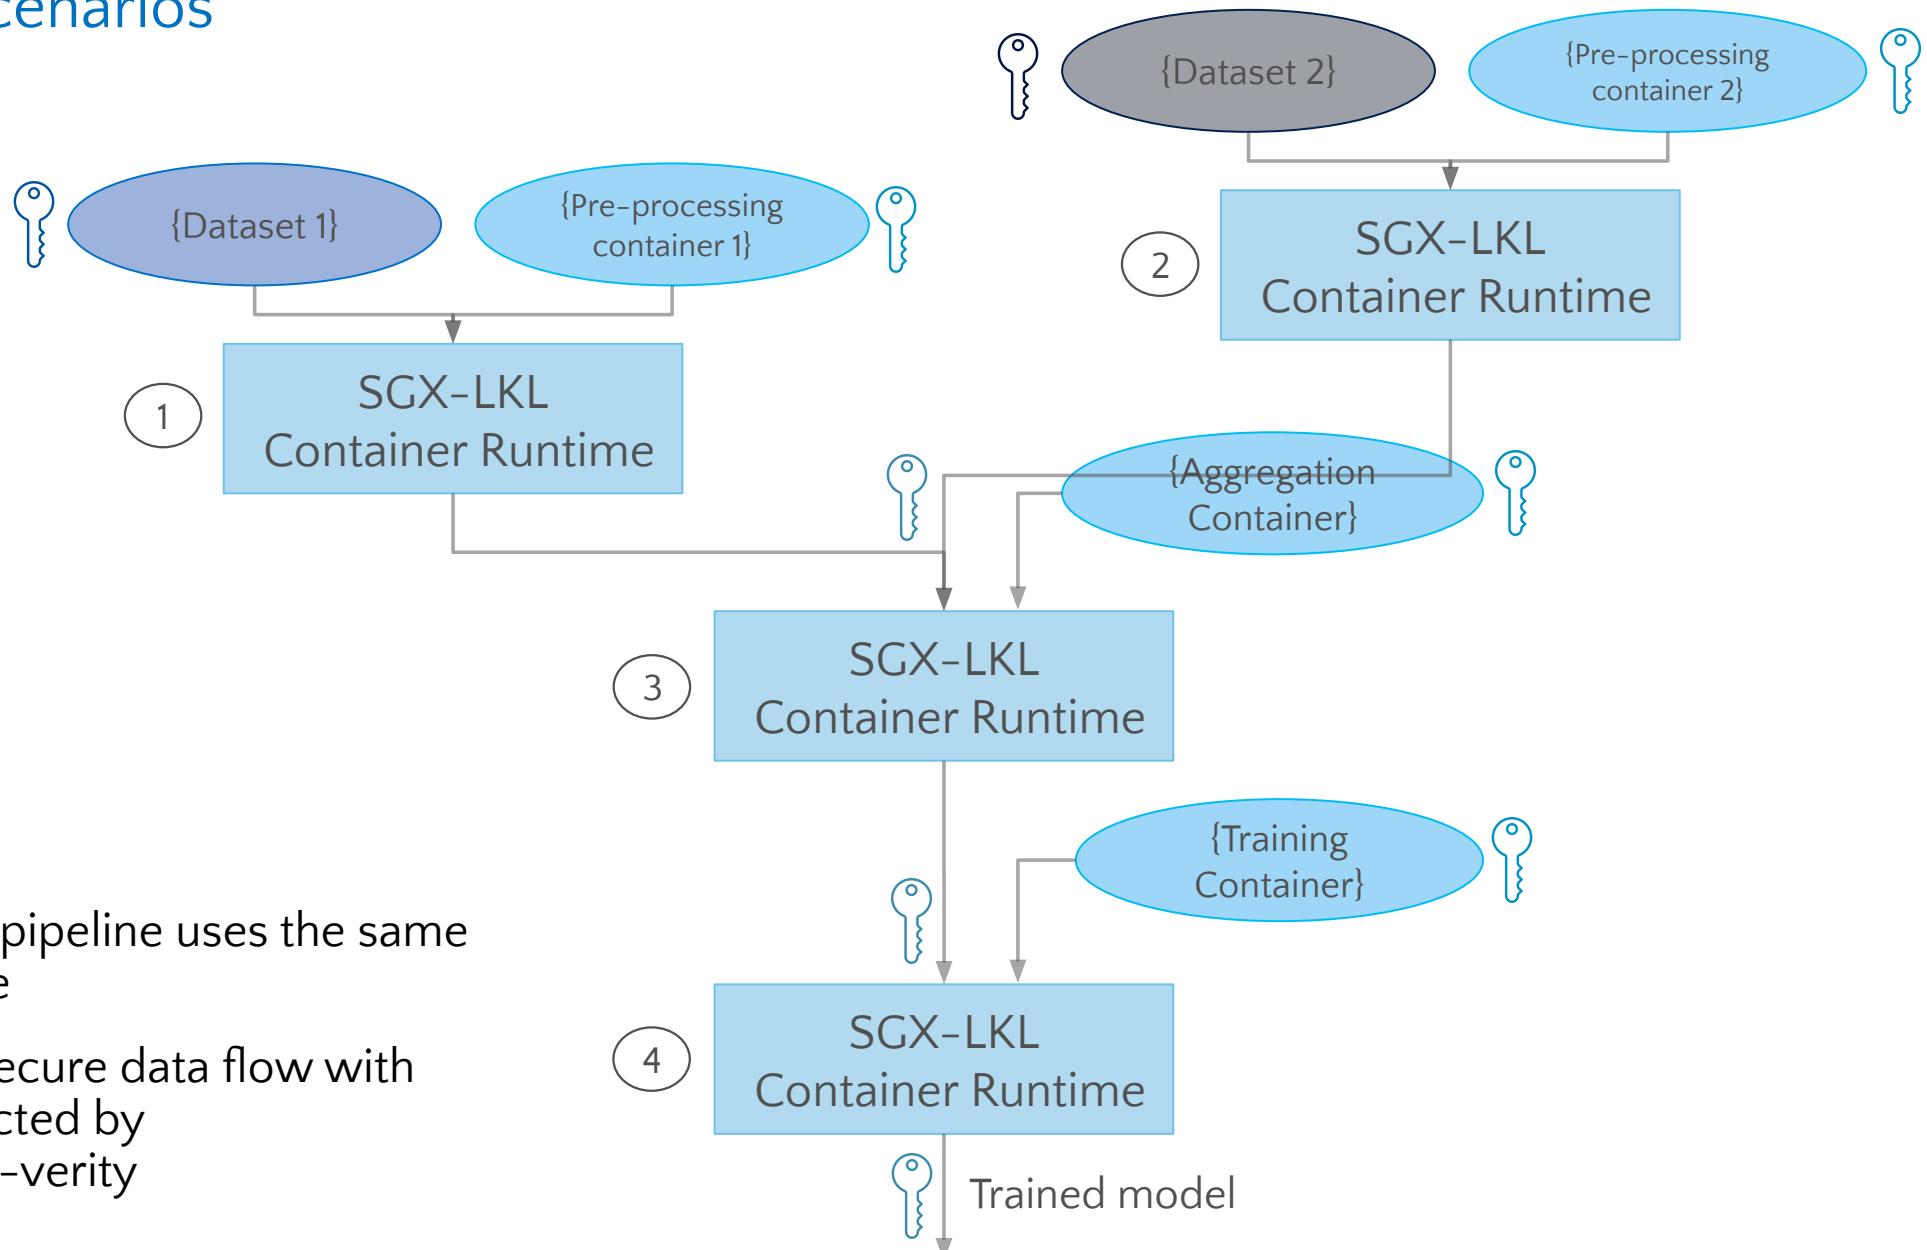
Each facility only sees their respective data sets (aka no one, not even cloud provider, can see all data or trained model, if necessary)

All facilities benefit from using trained model



# Confidential ML pipelines for Multi Party machine learning

## Healthcare scenarios



Every node of the pipeline uses the same LKL runtime image

Edges represent secure data flow with disk images protected by dm-crypt and dm-verity

# To go beyond...

## Azure confidential computing and Confidential ML links

Azure confidential computing solution page:

<https://azure.microsoft.com/en-us/solutions/confidential-compute/>

Azure confidential computing:

<https://azure.microsoft.com/en-us/blog/azure-confidential-computing/>

Ignite Mechanics Video with Mark Russinovich:

<https://youtu.be/Qu6sPOXDMU8>

Azure Confidential Compute (Virtual Machine):

<https://azuremarketplace.microsoft.com/en-us/marketplace/apps/microsoft-azure-compute.acc-virtual-machine-v2?tab=Overview>

Open Enclave SDK page:

<https://openenclave.io/sdk/>

Confidential ML:

<https://aka.ms/AIShow/ConfidentialML>

Responsible AI resources:

<https://www.microsoft.com/en-us/ai/responsible-ai-resources>

Azure responds to COVID-19:

<https://azure.microsoft.com/en-us/blog/azure-responds-to-covid19/>

To go beyond...



## Confidential Computing @ Microsoft Research (MSR)

### Reference material on Intel SGX use from

Multi-party machine learning:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/paper.pdf>

SQL Server with Haven:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/osdi2014-haven.pdf>

Map/reduce with VC3:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/vc3-oakland2015.pdf>

Preventing enclave information leaks:

<https://people.eecs.berkeley.edu/~rsinha/research/pubs/pldi2016.pdf>

Using side-channel page faults to extract JPG images:

<https://www.microsoft.com/en-us/research/wp-content/uploads/2017/06/atc17-final230.pdf>

PROTÉGER NOS DONNÉES LORS DE L'EXÉCUTION

# Homomorphic Encryption (HE)



# Homomorphic Encryption (HE)

## New type of encryption

Allows (un)limited computation on encrypted data

Enables outsourcing of data storage/processing, i.e. data enters / changes / leaves untrusted networks – always encrypted

- No information about the data is revealed to the untrusted network
- Results can only be revealed by data owner
- Think persistent cloud storage with computation (in Azure) : statistical computations, ML inference, secure search

Can be either public-key or private-key encryption

Post-quantum secure

## "Full homomorphic" encryption

$$D_K(C_K(n) + C_K(m)) = n + m$$

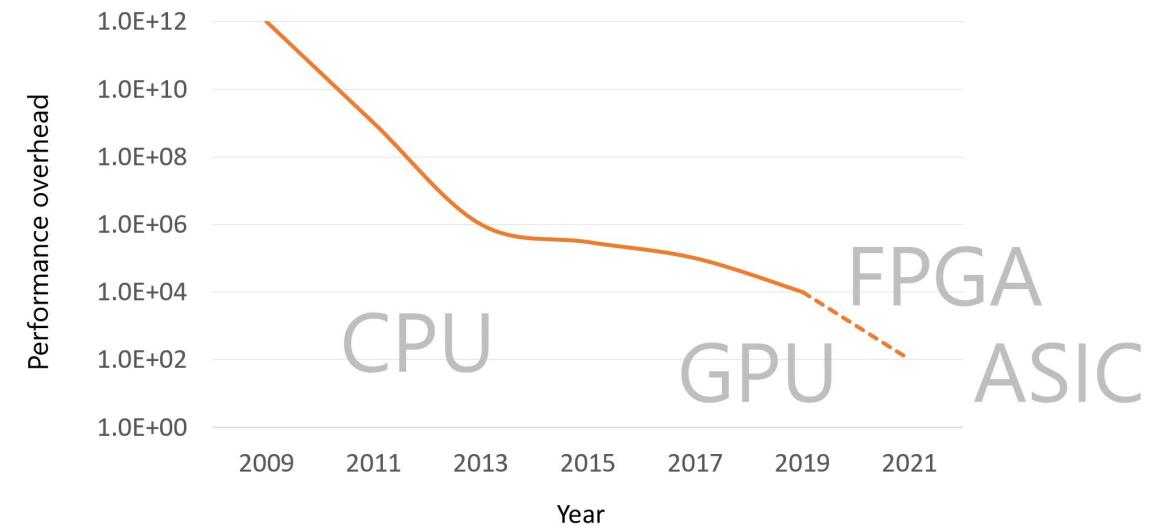
$$D_K(C_K(n) \times C_K(m)) = n \times m$$

## Some benefits

There is no hardware element in the TCB  
(so you don't have to trust hardware)

This type of encryption is applicable for some applications, but... its extra cost is still intolerable for general-purpose treatments

## Performance Overhead



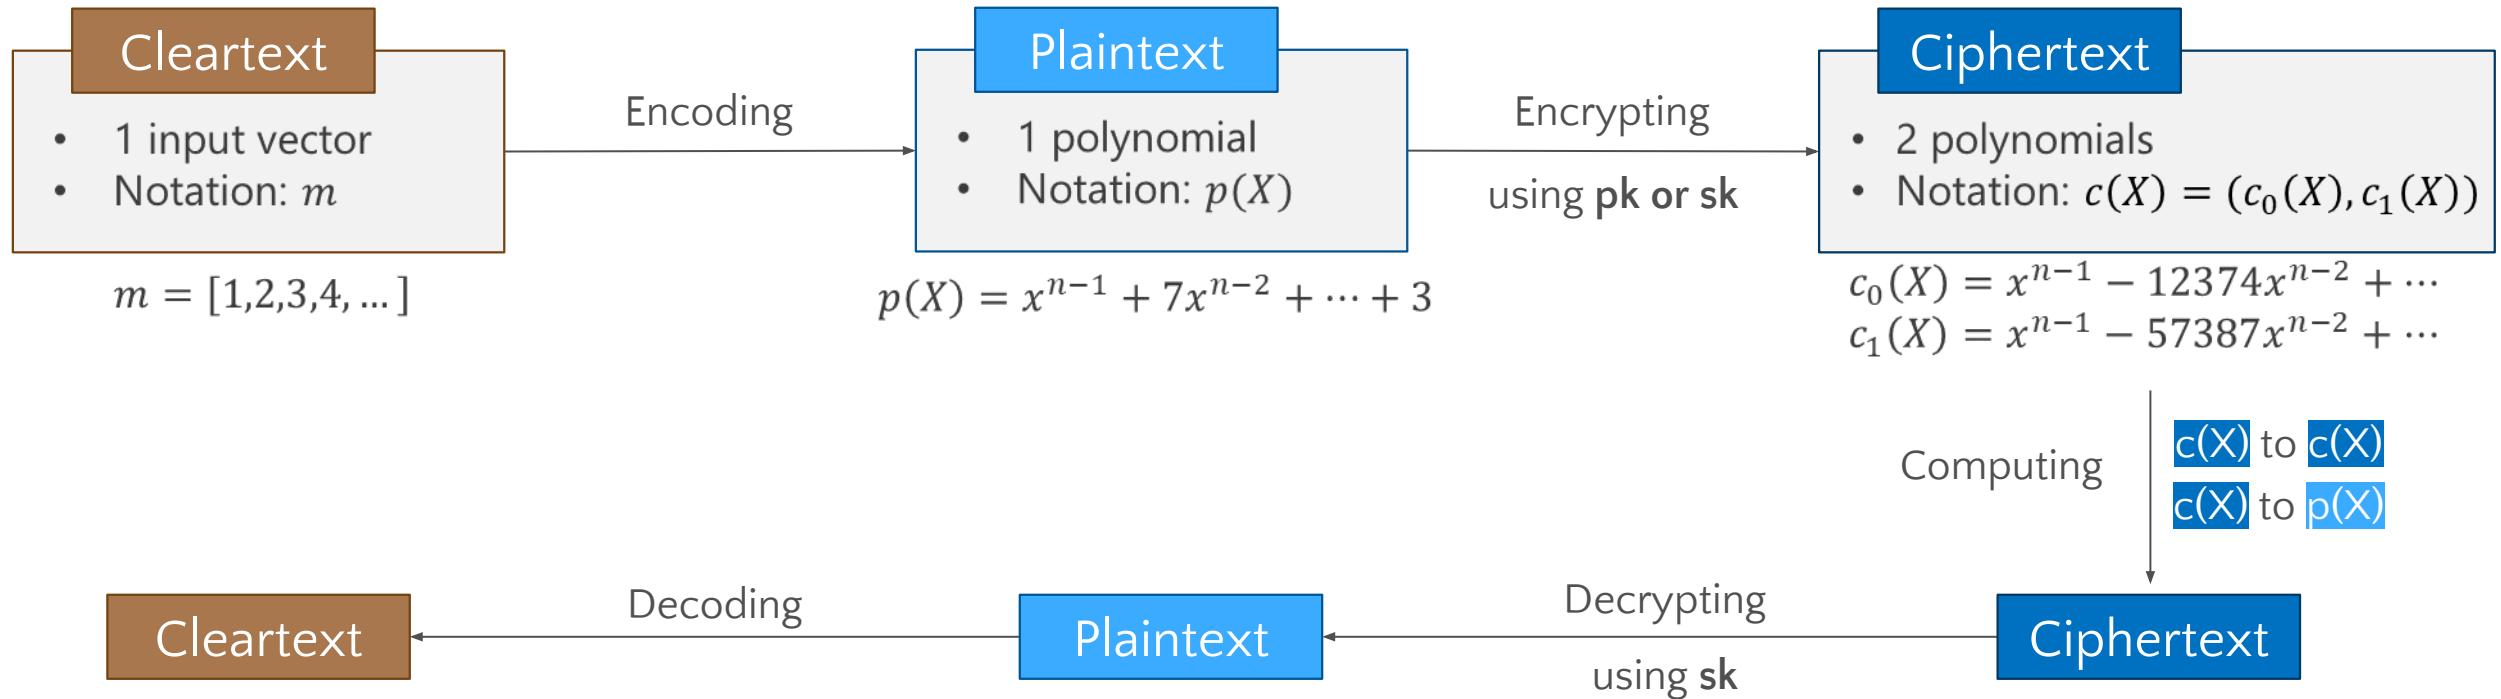
Invented in 2009 by Craig Gentry,  
@ Stanford university

Originally, 100 trillion times slower to make calculations on encrypted data than on unencrypted data

2 million times faster in 2013

# Homomorphic Encryption

## High level view



### Notes & Key points

- Homomorphic operations operate with polynomials example
- Microsoft SEAL can work both using
  - ✓ symmetric encryption (one secret key for encryption and decryption)
  - ✓ asymmetric encryption (one public key for encryption and one secret key for decryption)
- The encoding operation is also a homomorphism

# Homomorphic Encryption

## Encoding and decoding with SEAL: the CKKS encoder

### CKKS encoder

#### Principle

- Input: one vector of size  $n$
- Principle:
  - Encode by computing the Lagrangian polynomial interpolation on specific values (the roots of  $X^n + 1$ )
  - Decode by evaluating on the same values
- Vector view:  
 $m = [m_0, m_1, \dots, m_{\frac{n}{2}}, m_{\frac{n}{2}}] \rightarrow p(X) = [p_0, p_1, \dots, p_{n-1}]$

#### Example

$$m = [3 + 4i, 2 - i] \rightarrow p(X) = 45X^3 + 160X^2 + 91X + 160$$

#### Notes & Key points

#### About the CKKS encoder

- You can embed at most  $\frac{n}{2}$  numbers in one message
- $m$  can contain floats or complex numbers
- Coefficients of  $p(X)$  are computed modulo Q

#### Another view

- Considering:  $R = (r_0, r_1, \dots, r_{n-1})$  the roots of  $X^n + 1$
- Principle:

$$p(r_i) = \Delta m_i \text{ for } i \in \llbracket 0, \frac{n}{2} \rrbracket$$

- Vector view:

$$\text{with } p(X) = \begin{matrix} p_0 \\ p_1 \\ \dots \\ p_{\frac{n}{2}-1} \end{matrix}, \quad p(R) = \Delta m = \begin{matrix} \Delta m_0 \\ \Delta m_1 \\ \dots \\ \Delta m_{\frac{n}{2}-1} \end{matrix}$$

$$\text{because } p(R) = [p(r_0), p(r_1), \dots, p(r_{\frac{n}{2}-1})]$$

#### About all encoders:

- $n$  is always a power of 2
- $p(X)$  has always a degree of  $n-1$

# Homomorphic Encryption

## Encrypting and decrypting

### Step 1



#### Key generation

- Secret key:  $s_k$  (*randomly chosen*)
- Public key:  $p_k = (b, a)$ 
  - $a$  another sampled polynomial
  - $b = -as_k + e$
  - $e$  some random noise

### Step 2

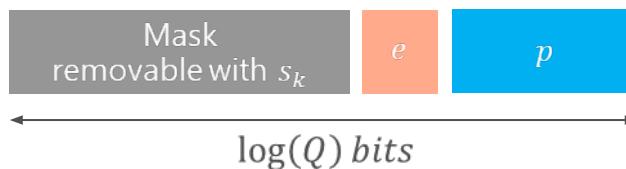


#### Encryption

Considering  $c = (c_0, c_1)$

- $c_0 = b + p = -as_k + e + p$
- $c_1 = a$

Thus, a ciphertext is structured as followed:



### Step 3



#### Decryption

Considering  $p$  as the encoded message  $m$ ,

$$\begin{cases} c_0(X) = -as_k(X) + p(X) \\ c_1(X) = a(X) \end{cases}$$

so

$$c_0 + c_1 s_k = -as_k + e + p + as_k$$

$$c_0 + c_1 s_k = e + p$$

### Notes & Key points

- Remember that  $s_k, a, b, c_0, c_1, p$  and  $e$  are **polynomials**
- **Decryption is not exact.** The requirement for a correct decryption is  $e \ll p$
- Each computation **increases the noise**. Once it reaches a certain point, correct decryption will be impossible.
- The **security** of the encryption relies on the hardness of finding  $s_k$  from  $b$  : this is called the *Ring Learning With Errors* problem, and it is **quantum secure**.

# Microsoft SEAL

## Computation : Native operations cheat sheet

### Scenarios



#### Addition

$c(X) \leftrightarrow c(X)$  or  $c(X) \leftrightarrow p(X)$

Performance cost:  $\vartheta(l \times n)$



#### Multiplication

$c(X) \leftrightarrow c(X)$  or  $c(X) \leftrightarrow p(X)$

Performance cost:  $\vartheta(l \times n \log n)$  in BFV,  
 $\vartheta(ln)$  in CKKS



#### Rescaling (CKKS)

Performed on  $c(X)$  only

Performance cost:  $\vartheta(l \times n \log n)$

### Foundation



#### Rotation

Performed on  $c(X)$  only

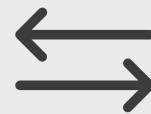
Performance cost:  $\vartheta(n \times \log n \log Q)$



#### Relinearization

Performed on  $c(X)$  only

Performance cost:  $\vartheta(l^2 \times n \log n)$



#### Modulus switching

Performed on  $c(X)$  only

Performance cost:  $\vartheta(l \times n)$  in BFV,  
 $\vartheta(l \times n \log n)$  in CKKS

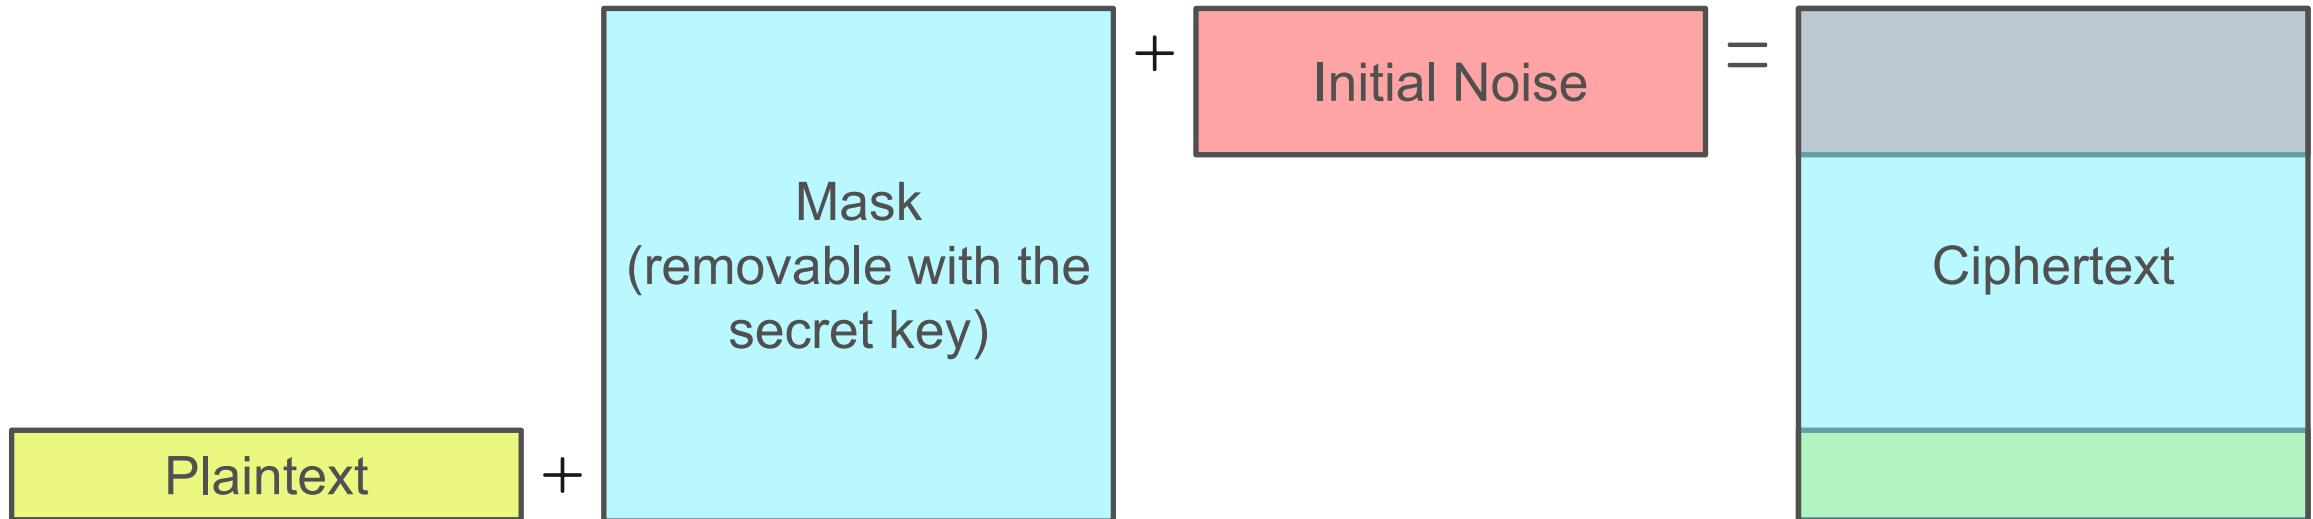
Note:  $l$  is the computational depth of the circuit

# Homomorphic Encryption (HE) considerations

## Anatomy of a ciphertext : Noise

The noise term is essential for homomorphic encryption schemes

Note also big message expansion



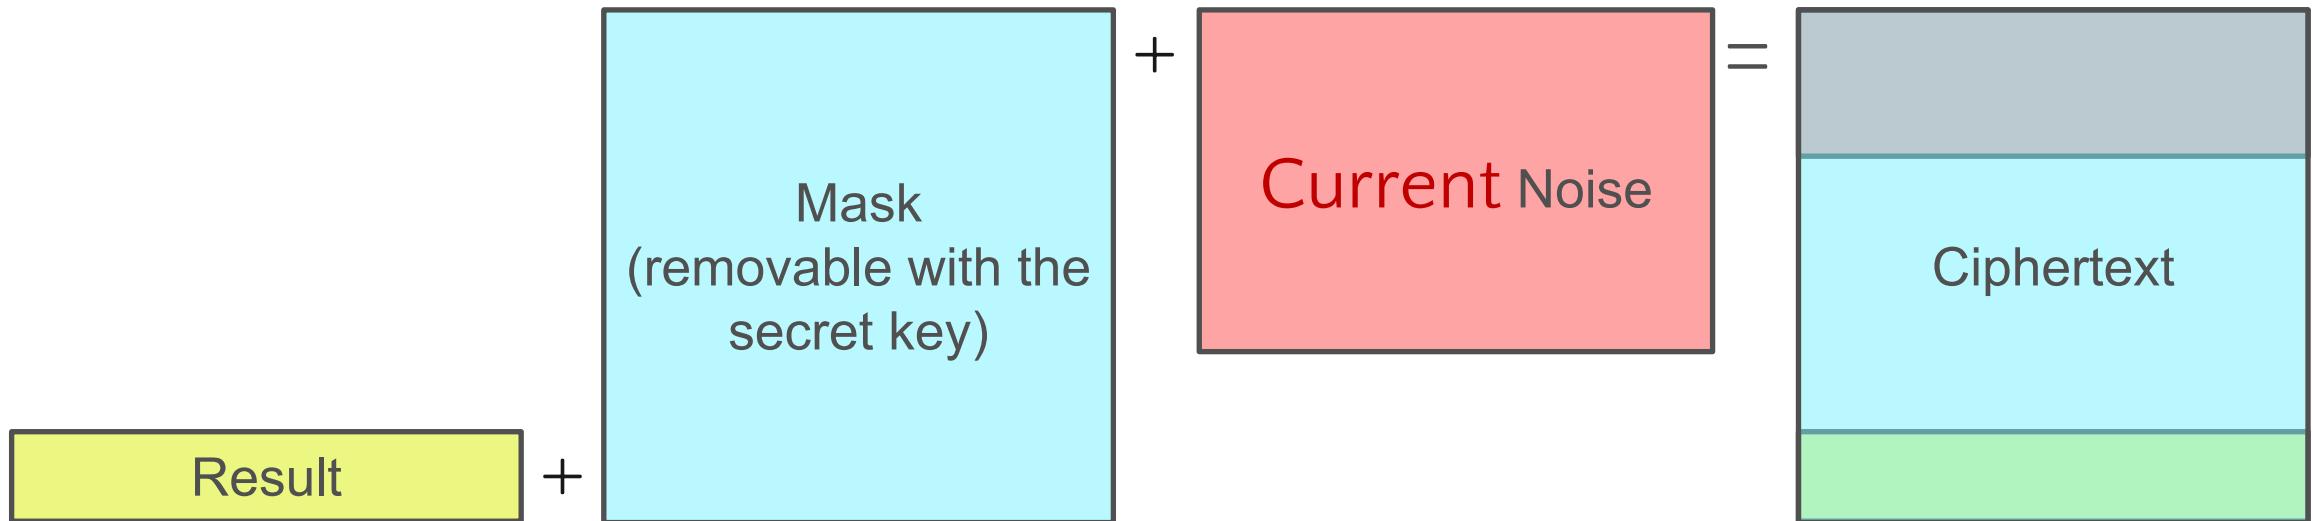
- Initial noise (of a fresh encryption) is small in terms of coefficients' size
- Independently of the initial noise, there is always a big message expansion : it is part of the security basis of this type of encryption

# Homomorphic Encryption (HE) considerations

## Anatomy of a ciphertext : after computation

After each level, noise increases

Note how message stays same size due to modular arithmetic

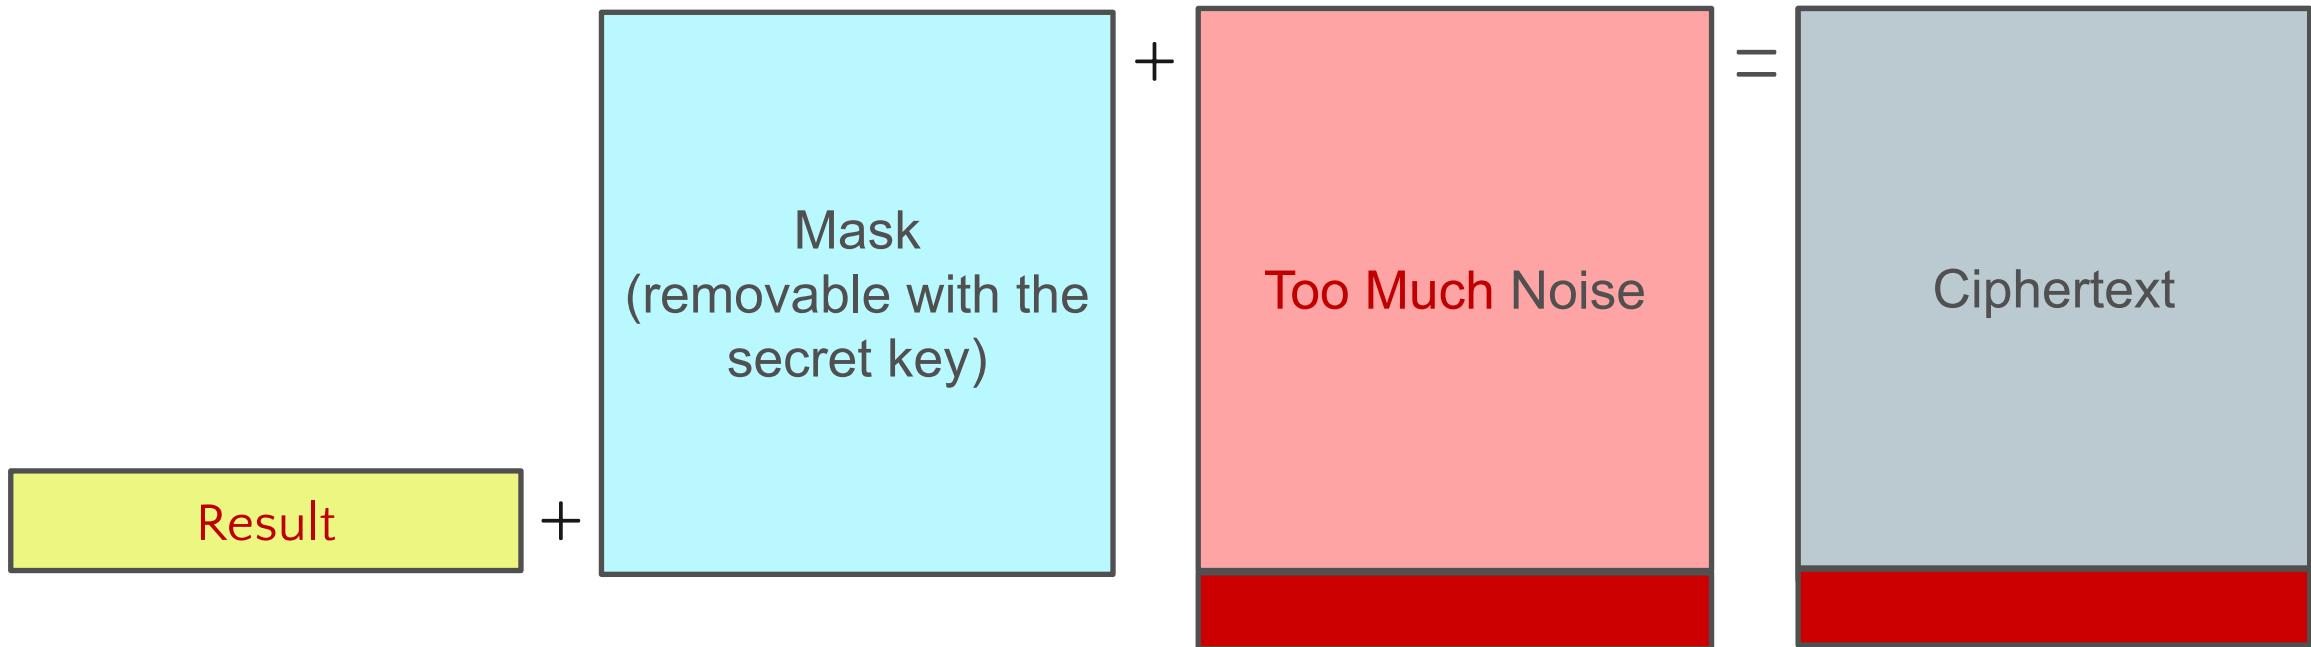


# Homomorphic Encryption (HE) considerations

## Anatomy of a ciphertext : Noise overflow

At some level, noise completely destroys the message

This is why only limited computation is possible



# Homomorphic Encryption (HE) considerations

## Anatomy of a ciphertext : decryption

Decrypting is removing the mask

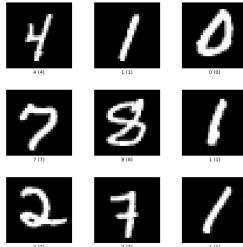
Note that there is a residual noise that must be managed



# Microsoft homomorphic encryption

## CryptoNets performance break-through

A breakthrough in 2016: evaluates neural net predictions on encrypted data



## Simple Encrypted Arithmetic Library (SEAL)

See next slide

Widely adopted by research teams worldwide

## Standardization in progress w/ community

<http://sealcrypto.org>

## Microsoft Open Sources Homomorphic Encryption Library “SEAL”

ED TARGETT EDITOR  
2ND JANUARY 2019

+ INCREASE / DECREASE TEXT SIZE -



Add to favorites



“Easy for both cryptography experts and novice practitioners”

Microsoft has open sourced a homomorphic encryption library developed by its [Cryptography Research group](#), saying it “strongly believes” the technology is ripe for use in real-world applications, as it makes the source code available on GitHub.

Homomorphic Encryption provides the unique ability to compute on data while it is encrypted – Microsoft has been leading a push to develop a standard for the emerging cryptographic technique, which has wide ranging industry potential.

Adoption of homomorphic encryption, for example, would allow users to process encrypted data in the cloud, without either downloading it for decryption on-premises, or providing a decryption key to a third party service provider.

Microsoft’s Simple Encrypted Arithmetic Library (SEAL) (open sourced early last month under an MIT Licence for free use) has already been

### MOST READ



Intel Names New CEO: Meet the New Boss...  
ED TARGETT 31ST JANUARY 2019



Airbus Hack: Aerospace and Security Giant Breached  
CONOR REYNOLDS 31ST JANUARY 2019



Five Things We Learned from Microsoft's Earnings  
ED TARGETT 31ST JANUARY 2019

### EDITOR'S PICK



Measuring DevOps Performance: The Key Indicators  
ED TARGETT 18TH JANUARY 2019



Four Ways Business Leaders Can Harness the \$500 Million Data Literacy Opportunity  
ED TARGETT 15TH JANUARY 2019



Creating a Winning Commerce Strategy with API-First Commerce  
ED TARGETT 14TH JANUARY 2019

# Microsoft SEAL

## Simple Encrypted Arithmetic Library

A state-of-the-art homomorphic encryption library  
from Microsoft Research (MSR)

Encryption and decryption: Low-level primitives for encrypted arithmetic (add, subtract, multiply)

GPU/FPGA acceleration (up to +50x perf, resp. up to +120x perf)

SEAL 1.0 released under MSR-LA in 2015

SEAL 3.5.0 released under MIT license in April 2020

<https://github.com/microsoft/SEAL/blob/master/Changes.md>

Written in C++ 17; includes .NET Standard wrapper for public API (> SEAL 3.2)

- Supporting Azure Functions : serverless Secure Compute

From open source community:

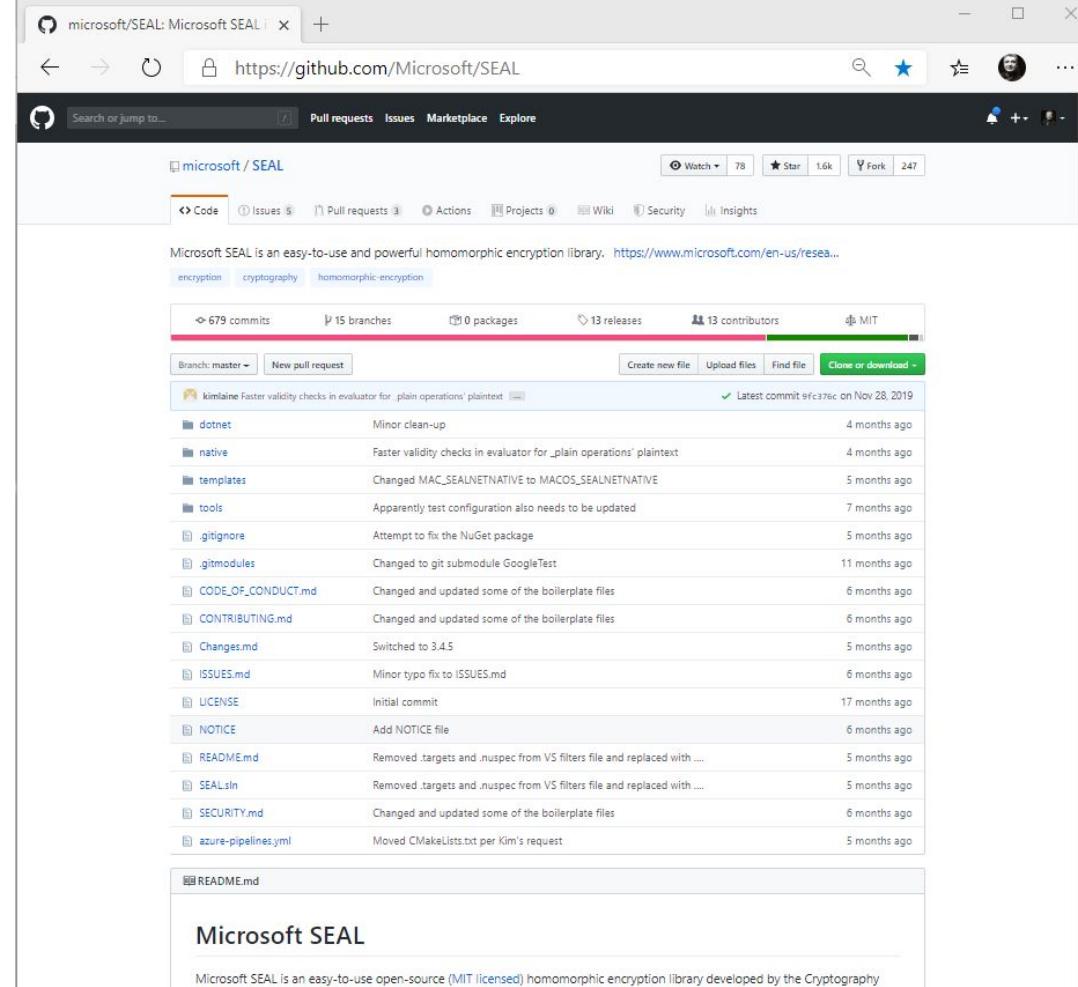
- [PyHeal](#) (Python wrappers from Accenture), [SEAL-python](#) (yet another wrapper for SEAL 3.4.5)
- [node-seal](#) (JavaScript wrappers)
- [nGraph HE Transformer](#) (from Intel)
- TenSEAL (from OpenMinded)

Support to:

- Operating System: Windows, Linux, macOS, FreeBSD
- Smartphones : Android

Samples' repos on GitHub:

- [SEAL-Demo](#)
- [algorithms-in-SEAL](#)
- [bootcamp](#)
- Etc.



# Applications of homomorphic encryption

## Implementing Private AI

Implementing neural networks directly against SEAL is very challenging...

CHET is an optimizing compiler for fully-homomorphic neural-network inferencing

ONNX Runtime is Microsoft's open-source inference engine for ONNX, see [RendezVous S03R03](#) "Open Neural Network eXchange (ONNX) De quoi s'agit-il ?"

An integration into ONNX Runtime allows our Private AI efforts to benefit from the large ecosystem around ONNX, which implies to expose capabilities for encrypted inferencing in ONNX Runtime:

- Adding a CHET execution provider
- Extending handling of encrypted inputs and outputs in ONNX Runtime

For models that CHET provides kernels for, compatible ONNX models and data will be provided.

EVA is a generic low-level language for homomorphic encryption

CHET compiles neural networks into EVA programs

EVA programs are optimized and compiled into SEAL API calls

SECLUD is a PaaS offering that will offer powerful but easy-to-understand privacy-preserving services through Azure

Private data lookup, statistics, data collaboration, and ML predictions

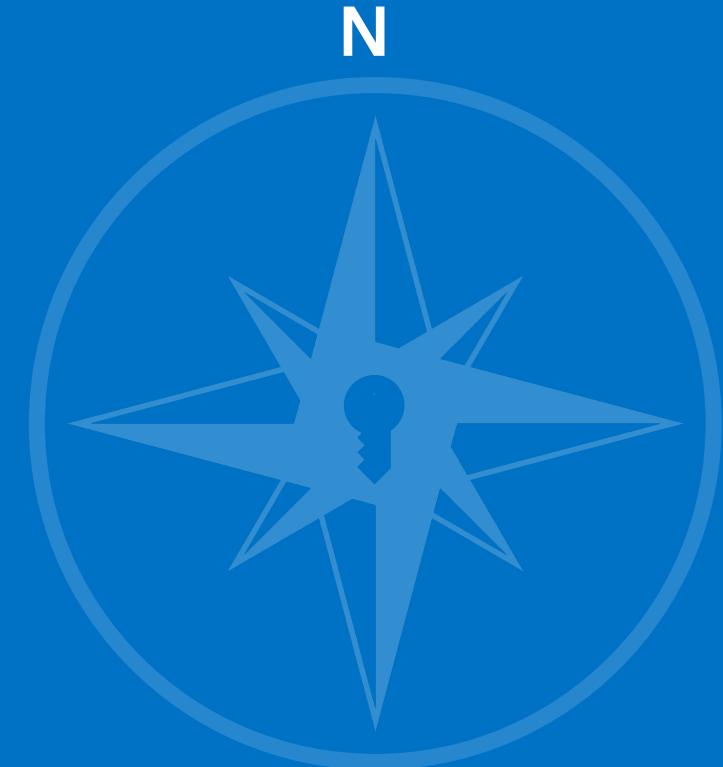
One core area of investment is enabling CHET-based inferencing in SECLUD, and thus integrating encrypted ONNX Runtime inferencing

# Demo

First steps with TenSEAL

Code available at:

<https://github.com/OpenMined/TenSEAL/tree/master/tutorials>



# Demo

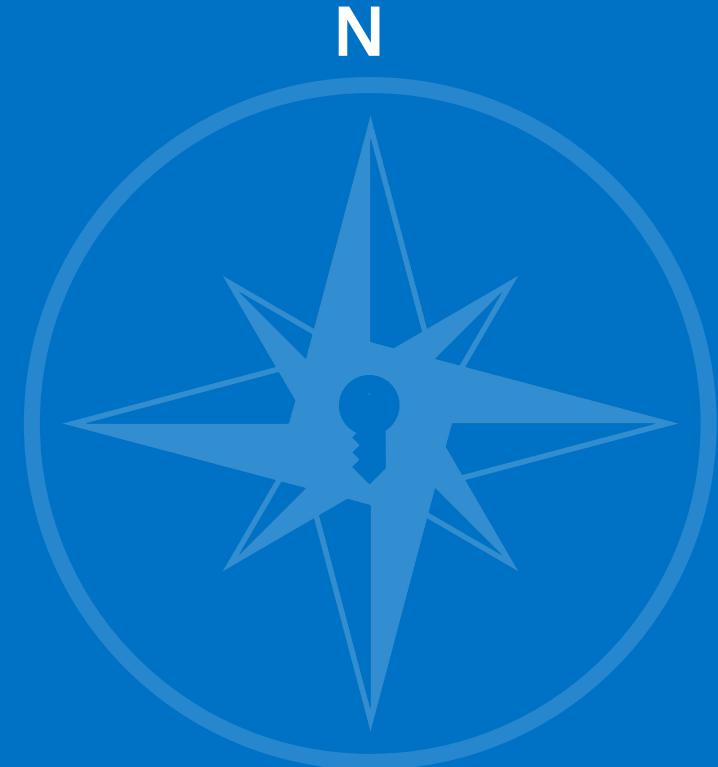
## Homomorphic Random Forests

Paper available at:

<https://arxiv.org/pdf/2006.08299.pdf>

Code available at:

<https://github.com/dhuynh95/cryptotree>



# To go beyond...

On [docs.microsoft.com](https://docs.microsoft.com)

[How to deploy an encrypted inferencing web service](#)

- See [Tutorial #3: Deploy an image classification model for encrypted inferencing in Azure Container Instance \(ACI\)](#)

[Medium series of articles on Homomorphic Encryption intro](#)

Series of article on HE from the ground up, with a Python implementation from scratch of CKKS:

- [Overview and use cases](#)
- [HE landscape and CKKS](#)
- [Encoding and decoding in CKKS](#)

To go beyond...



## Private AI Bootcamp December 2019

With a series of videos on:

- Intro to Homomorphic Encryption,
- Microsoft SEAL (+ [bootcamp](#) repo on GitHub),
- Intro to CKKS,
- Techniques in PPML (+ [algorithms-in-SEAL](#) repo on GitHub),
- Building Applications with Microsoft SEAL,
- Etc.

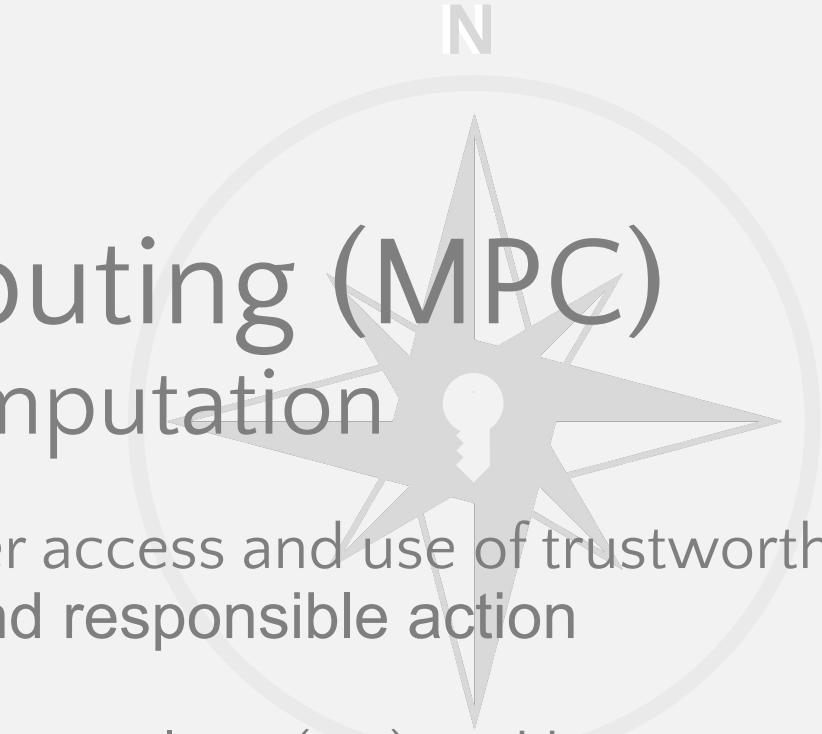
## Private AI February 2020

# Secure Multi-Party Computing (MPC)

## Enabling secure collaborative computation

Solving today's complex challenges require greater access and use of trustworthy data to create new insights and facilitate **responsive and responsible action**

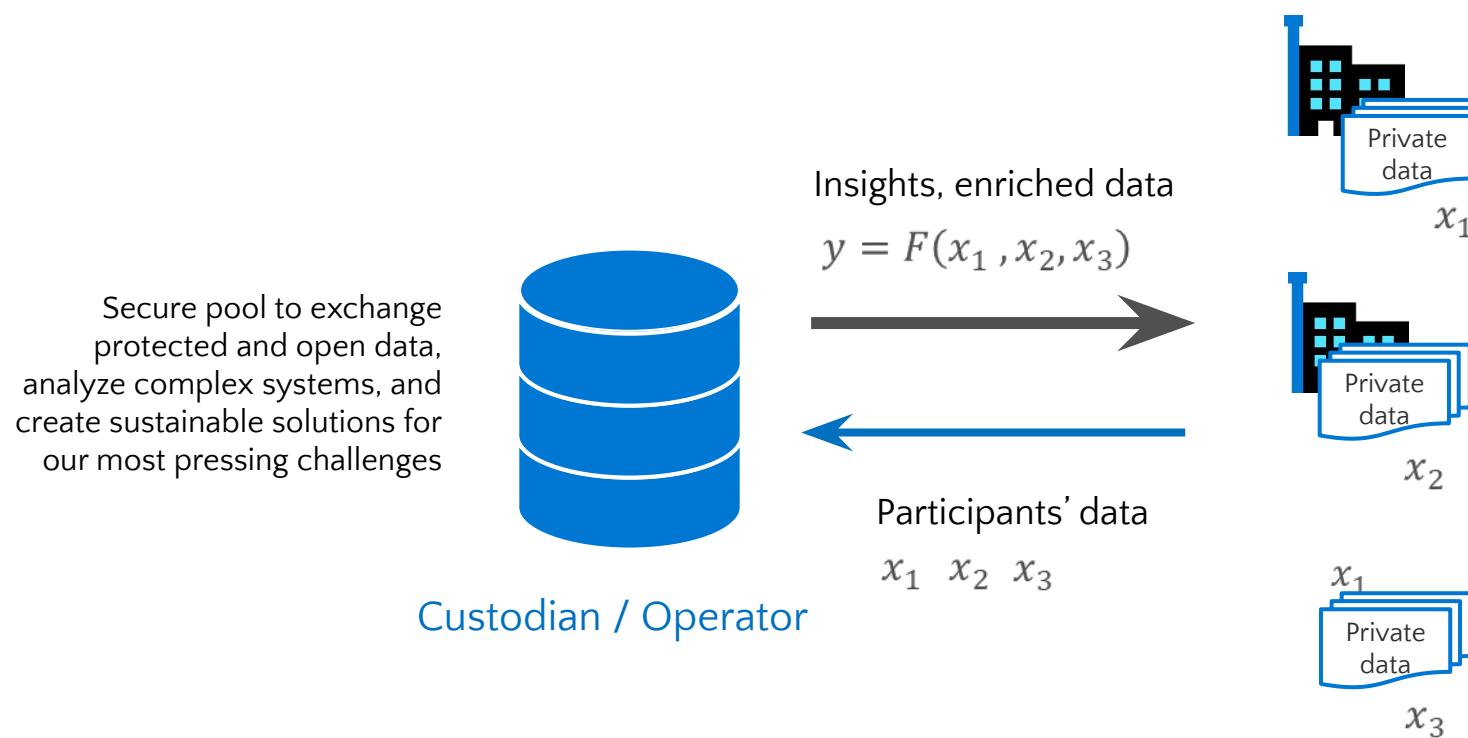
MPC has different trade-offs from trusted hardware enclave (CC) and homomorphic encryption (HE)



# Secure Multi-Party Computation (MPC) pattern

**A custodian provides a secure pool that meets requirements for confidentiality and compliance**

Input privacy:  $x_i$  is not leaked to parties  $\neq i$



# Principle

1 Connect workers



MODEL



2 Build shares



MODEL



3 Exchange shares



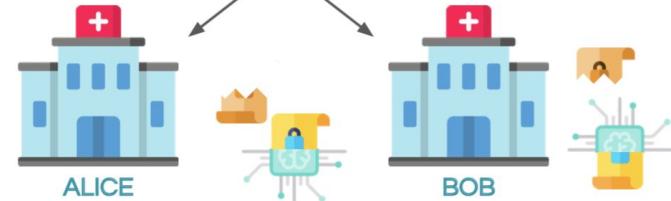
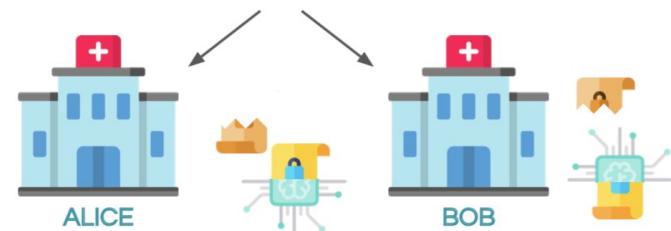
MODEL



4 Send encrypted model to learn on



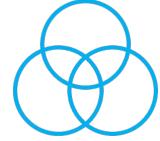
MODEL



Source: PySyft from OpenMined

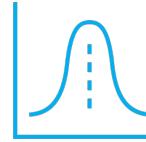
<https://github.com/OpenMined/PySyft>

# Secure MPC principles



## Additive shares

- If  $N$  people participate in a computation, data is sharded into  $N$  shares such that knowing up to  $N-1$  shares provides no information on the input, but only by summing all  $N$  shares we can reconstruct the data



## Randomized shares

- To provide privacy for each party when data is shared, we sometime need to have a third party providing random masks  $e$  to each party so that they can disclose  $y = x + e$ , instead of just  $x$  to hide the initial input



## Secure computations

- Using shares and random mask, complex operations can be performed such as matrix multiplication, comparisons, MaxPooling, etc.

# SMPC principles

## Example of multiplication

- Have  $\langle x \rangle, \langle y \rangle$ , want  $\langle x \cdot y \rangle$ .
- Use random triple  $\langle a \rangle, \langle b \rangle, \langle a \cdot b \rangle$
- Compute and open  $\langle x + a \rangle, \langle y + b \rangle$
- Observe:

$$\begin{aligned} x \cdot y &= (x + a - a) \cdot (y + b - b) \\ &= (x + a) \cdot (y + b) - (x + a) \cdot b - a \cdot (y + b) + a \cdot b \end{aligned}$$

The diagram illustrates the components of the multiplication equation. Blue boxes represent 'opened' components ( $x + a$ ,  $y + b$ ,  $a$ ,  $b$ ) and red boxes represent 'preprocessed' components ( $x$ ,  $y$ ,  $a \cdot b$ ). Arrows point from each term in the equation to its corresponding colored box.

# Secure MPC protocols

EzPC project from Microsoft Research India

## SecureNN: Efficient and Private Neural Network Training (2019)

Source: <https://www.microsoft.com/en-us/research/uploads/prod/2018/09/securenn.pdf>

SecureNN is a 3-party SMPC framework with various building blocks such as matrix multiplication, ReLU, convolution, MaxPool, etc. for the training and inference of Neural Networks such that no party learns anything about other people's data

## CrypTFlow: Secure TensorFlow Inference (2020)

Source: <https://www.microsoft.com/en-us/research/uploads/prod/2019/09/CrypTFlow.pdf>

CrypTFlow is a Secure MPC system with three components:

1. Athos: a TensorFlow end-to-end compiler for SMPC backends
2. Porthos: a state of the art semi-honest SMPC protocol for TensorFlow inference
3. Aramis: a technique to convert semi-honest Secure MPC protocols to MPC protocols with malicious security using hardware integrity solutions such as Intel SGX (see previous Confidential Computing part)

# Applications of secure MPC

Auction, Voting, Social research

Distributed signing / Key management

Private set intersection / Location based services

How can two companies find which customers they have in common without revealing their customers to each other?

Collaborative Machine Learning on private data

Huge amount of internal and external interest; high-profile project

(Your favorite computing scenario, with privacy added :-))

# As a quick conclusion of this Common Talk

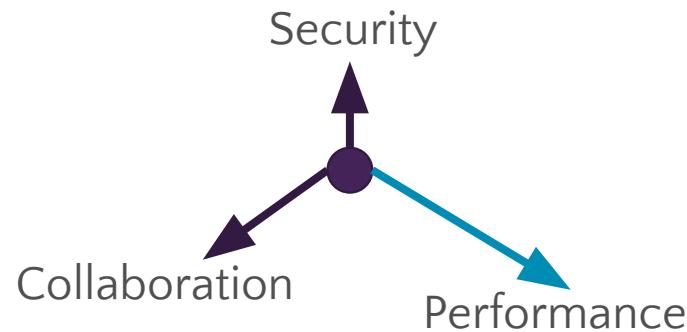
A review and comparison of our three main (mitigation) technologies' topics

	Confidential Computing (CC) (w/ secure hardware)	Homomorphic Encryption (HE)	Secure Multi-Party Computation (MPC)
Performance	-	Compute-bound	Network-bound
Privacy	Trusted Hardware	Encryption	Encryption / Non-collusion
Non-interactive	✓	✓	✗
Cryptographic security	✗ (known attacks)	✓	✓
Pros & Cons	Different Security Model Side-Channel Attacks	High Computation Overhead	High Communication High Number of Interactions

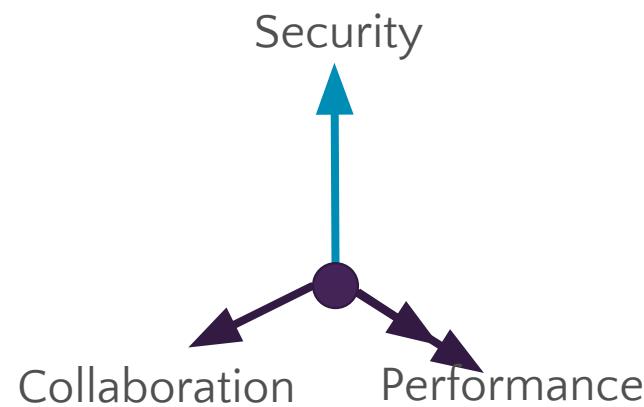
# As a conclusion of this Common Talk

A review and comparison of our three main technologies' topics, see also the wrap-up of the workshop for additional considerations

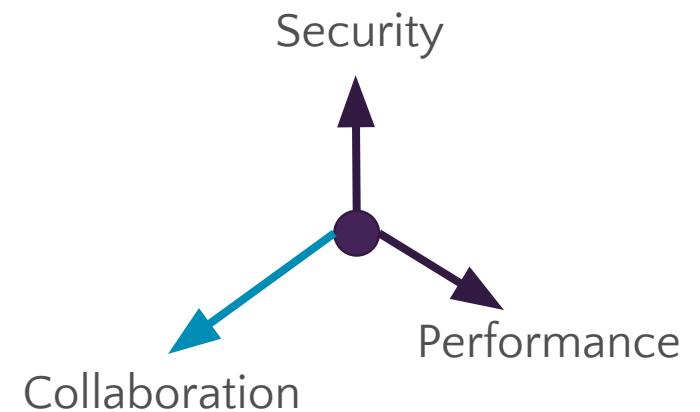
Confidential Computing (CC)  
w/ hardware enclaves



Homomorphic Encryption  
(HE)



Secure Multi Party  
Computation (MPC)



# Q&A

N



# Thanks!

ευχαριστώ Salamat Po متشرّم شكرًا Grazie

благодаря ありがとうございます Kiitos Teşekkürler 谢谢

ឧបរក្សាទុលេខ obrigado شكريه Terima Kasih Dziękuję

Hvala Köszönöm Tak Dank u Wel **дякую** Tack

Mulțumesc спасибо Danke Cám ơn Gracias

多謝晒 Říkám הַתִּל දෙන්ගි Děkuji 감사합니다

[https://aka.ms/DataInUseProtection  
WS](https://aka.ms/DataInUseProtectionWS)



Attribution 4.0 International  
(CC BY 4.0)

This presentation is shared under the license "Attribution 4.0 International (CC BY 4.0)". To see a copy of this license, visit <https://creativecommons.org/licenses/by/4.0/legalcode> or write to Creative Commons, PO Box 1866, Mountain View, CA 94042 USA

# Differential Privacy (DP)

## Protecting privacy while using data



# Differential Privacy (DP)

There is a significant amount of data that is inaccessible because of privacy

For most scenarios, data anonymization doesn't work

“Anonymized data isn’t”

- Cynthia Dwork, Microsoft Research (MSR)

Data scientists, analysts, scientific researchers and policy makers often do not need an individual's data

All they need is the signal

Differential privacy provides a way to get insights from the data, but without affecting the privacy of the individuals

A technique that offers strong privacy assurances, preventing data leaks and re-identification of individuals in a dataset

The screenshot shows the first page of a Microsoft Word document titled "Differential Privacy for Everyone". The Microsoft logo is at the top left. The title is centered above a section titled "Abstract". Below the abstract is a detailed introduction to differential privacy, mentioning its formal definition and applications. At the bottom of the page is a copyright notice and a page number "1".

Differential Privacy:  
A Primer for a Non-technical Audience\*  
(Preliminary version)

Kobbi Nissim<sup>1</sup>, Thomas Steinke<sup>2</sup>, Alexandra Wood<sup>3</sup>, Mark Bun<sup>2</sup>, Marco Gaboardi<sup>4</sup>, David R. O'Brien<sup>3</sup>, and Salil Vadhan<sup>2</sup>

<sup>1</sup>Department of Computer Science, Georgetown University.  
kobbi.nissim@georgetown.edu.

<sup>2</sup>Center for Research on Computation and Society, Harvard University.  
{tsteinke|mbum|salil}@seas.harvard.edu.

<sup>3</sup>Berkman Klein Center for Internet & Society, Harvard University.  
{awood|dobrien}@cyber.law.harvard.edu.

<sup>4</sup>University at Buffalo, The State University of New York. gaboardi@buffalo.edu.

March 3, 2017

## Abstract

This document is a primer on *differential privacy*, which is a formal mathematical framework for guaranteeing privacy when analyzing or releasing statistical data. It recently emerged from the theoretical computer science literature, differential privacy is now in its initial stages of implementation and use in various academic, industry, and government settings. Using intuitive illustrations and limited mathematical formalism, this document provides an introduction to differential privacy for non-technical practitioners, who are increasingly tasked with making decisions with respect to differential privacy as it grows more widespread in use. In particular, the examples in this document illustrate ways in which social scientists can conceptualize the guarantees provided by differential privacy with respect to the decisions they make when managing personal data about research subjects and informing them about the privacy protection they will be afforded.

**Keywords:** differential privacy, data privacy, social science research

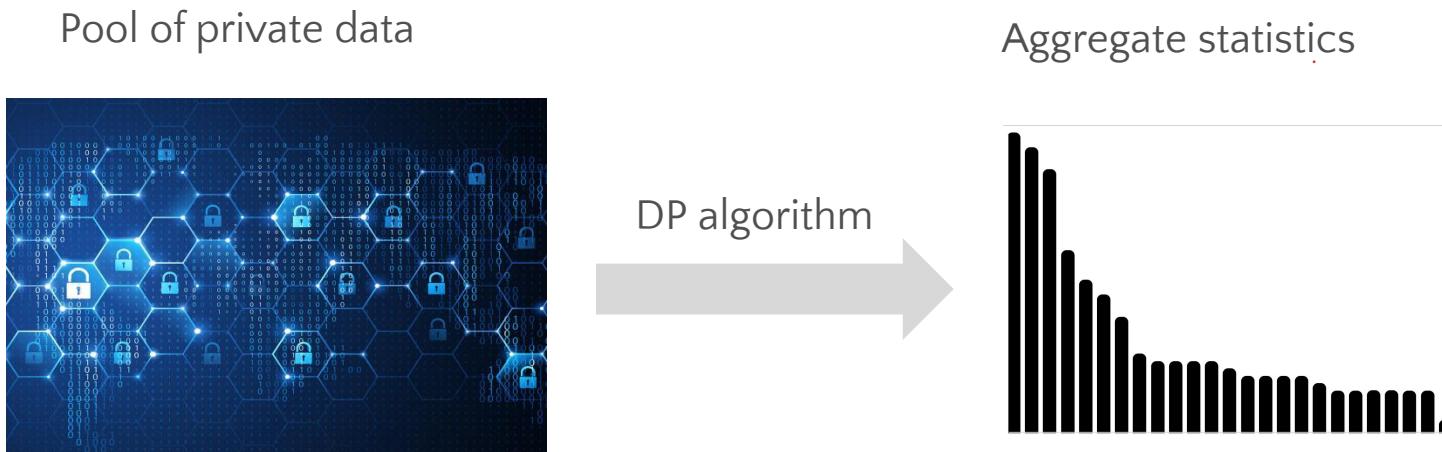
\*This document is the product of a working group of the *Privacy Tools for Sharing Research Data* project at Harvard University (<http://privacystools.seas.harvard.edu>). The working group discussions were led by Kobbi Nissim, Kobbi Nissim, Thomas Steinke, and Alexandra Wood were the lead authors of this document. Working group members include Caper Gooden, Deborah Hurley, Georgia Kellaris, Daniel Muisse, and Michel Reymond for their many valuable comments on earlier versions of this document. This material is based upon work supported by the National Science Foundation under Grant No. 1237235, as well as by the Alfred P. Sloan Foundation.

<sup>1</sup>Work towards this document was completed while the author was at the Center for Research on Computation and Society at Harvard University.

# Differential Privacy (DP)

A class of algorithms that facilitate computing and releasing aggregate statistics from sensitive (private) data, while ensuring that privacy is not compromised

Large pool of sensitive data, noisy aggregates are still valuable



DP statistics:

- Noisy. Noise conceals contributions of individual users
- Can be safely handed to an all-powerful adversary, with no risk of a privacy leak

Introduced by [Cynthia Dwork](#) and her collaborators at MSR in 2006

See [Differential Privacy](#)

Gradually came to dominate the area of privacy research

- Many thousands of citations to Cynthia et al. classical papers
- 2017 Gödel prize
- Deployments by major IT companies, US Census Bureau

Two main flavors

1. Local Differential Privacy (telemetry): simple statistics
2. Global Differential Privacy (trusted data collector): more complex aggregates including Machine Learning models

# Local Differential Privacy

## Setting

Data collector wants to understand global trends in the users' data

Data collector is not trusted, cannot collect raw data

## Local DP algorithm

Users carefully randomize their private data before sending it to data collector

Randomization provides plausible deniability

Noise becomes immaterial when data is collected from many users. Global trends are exposed

## Usage in telemetry

Facilitates collection of data that cannot be collected in the raw form

Suitable to collect basic statistical aggregates (mean, histogram, frequent items, etc.)

E.g. Telemetry in Windows: see [Collecting Telemetry Data Privately](#)

# Global Differential Privacy

## Setting

Database curator has access to a pool of private users' data

Needs to disclose some aggregate statistics to an analyst

## Global DP algorithms

Database curator gives noisy responses to analyst's queries: if the analyst wants  $f(D)$ , he gets  $f(D)+\text{noise}$

Noise is random, and hides a contribution of a single user

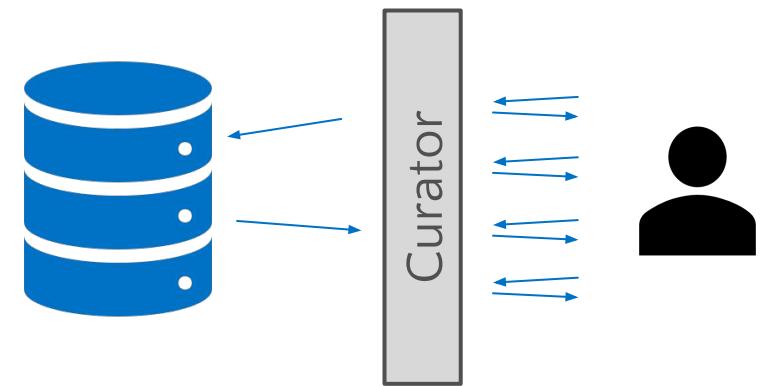
Any specific output is roughly as likely to be produced whether or not a single record is removed from the DB

## Usage examples

Office: expose the set of frequent e-mail responses, see [Assistive AI Makes Replying Easier](#)

LinkedIn (Ad campaign statistics), see [LinkedIn's Audience Engagements API: A Privacy Preserving Data Analytics System at Scale](#)

US Census Bureau: every published aggregate statistics for 2020 US Census will use DP

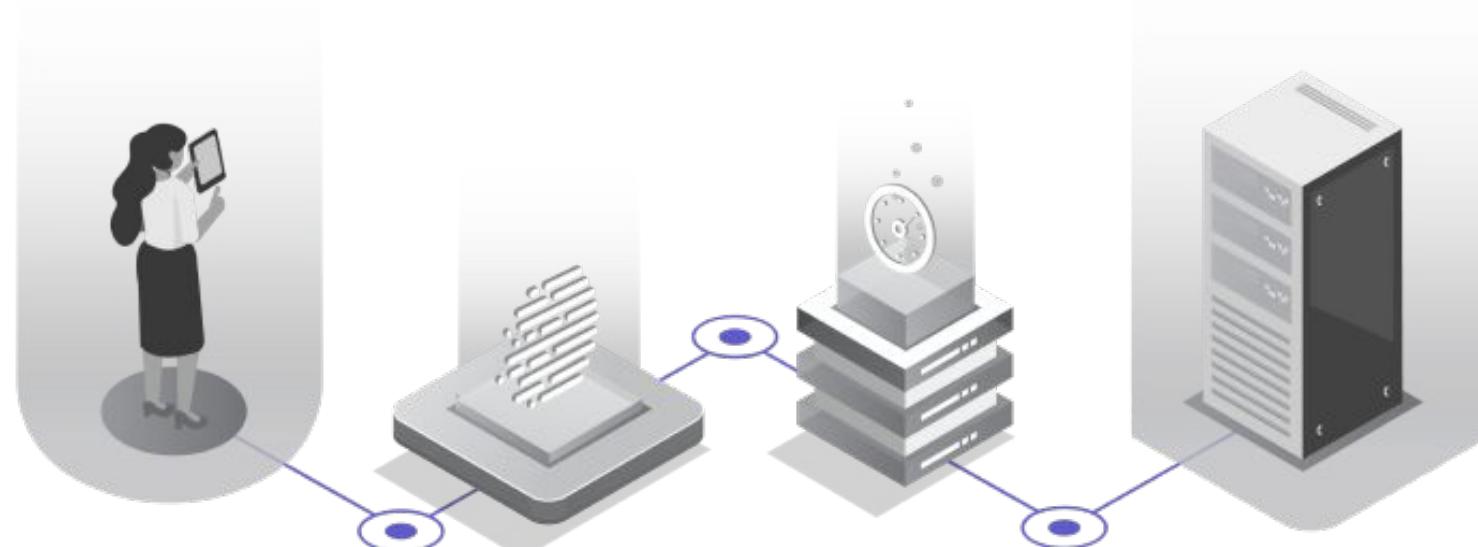


# Differential Privacy for Machine Learning and Analytics

## Introducing WhiteNoise

An open source project jointly developed by Microsoft, Harvard's Institute for Quantitative Social Science (IOSS) and the School of Engineering and Applied Sciences (SEAS) and

Currently supports scenarios where the data user is trusted by the data owner and wants to ensure that their releases or publications do not compromise privacy



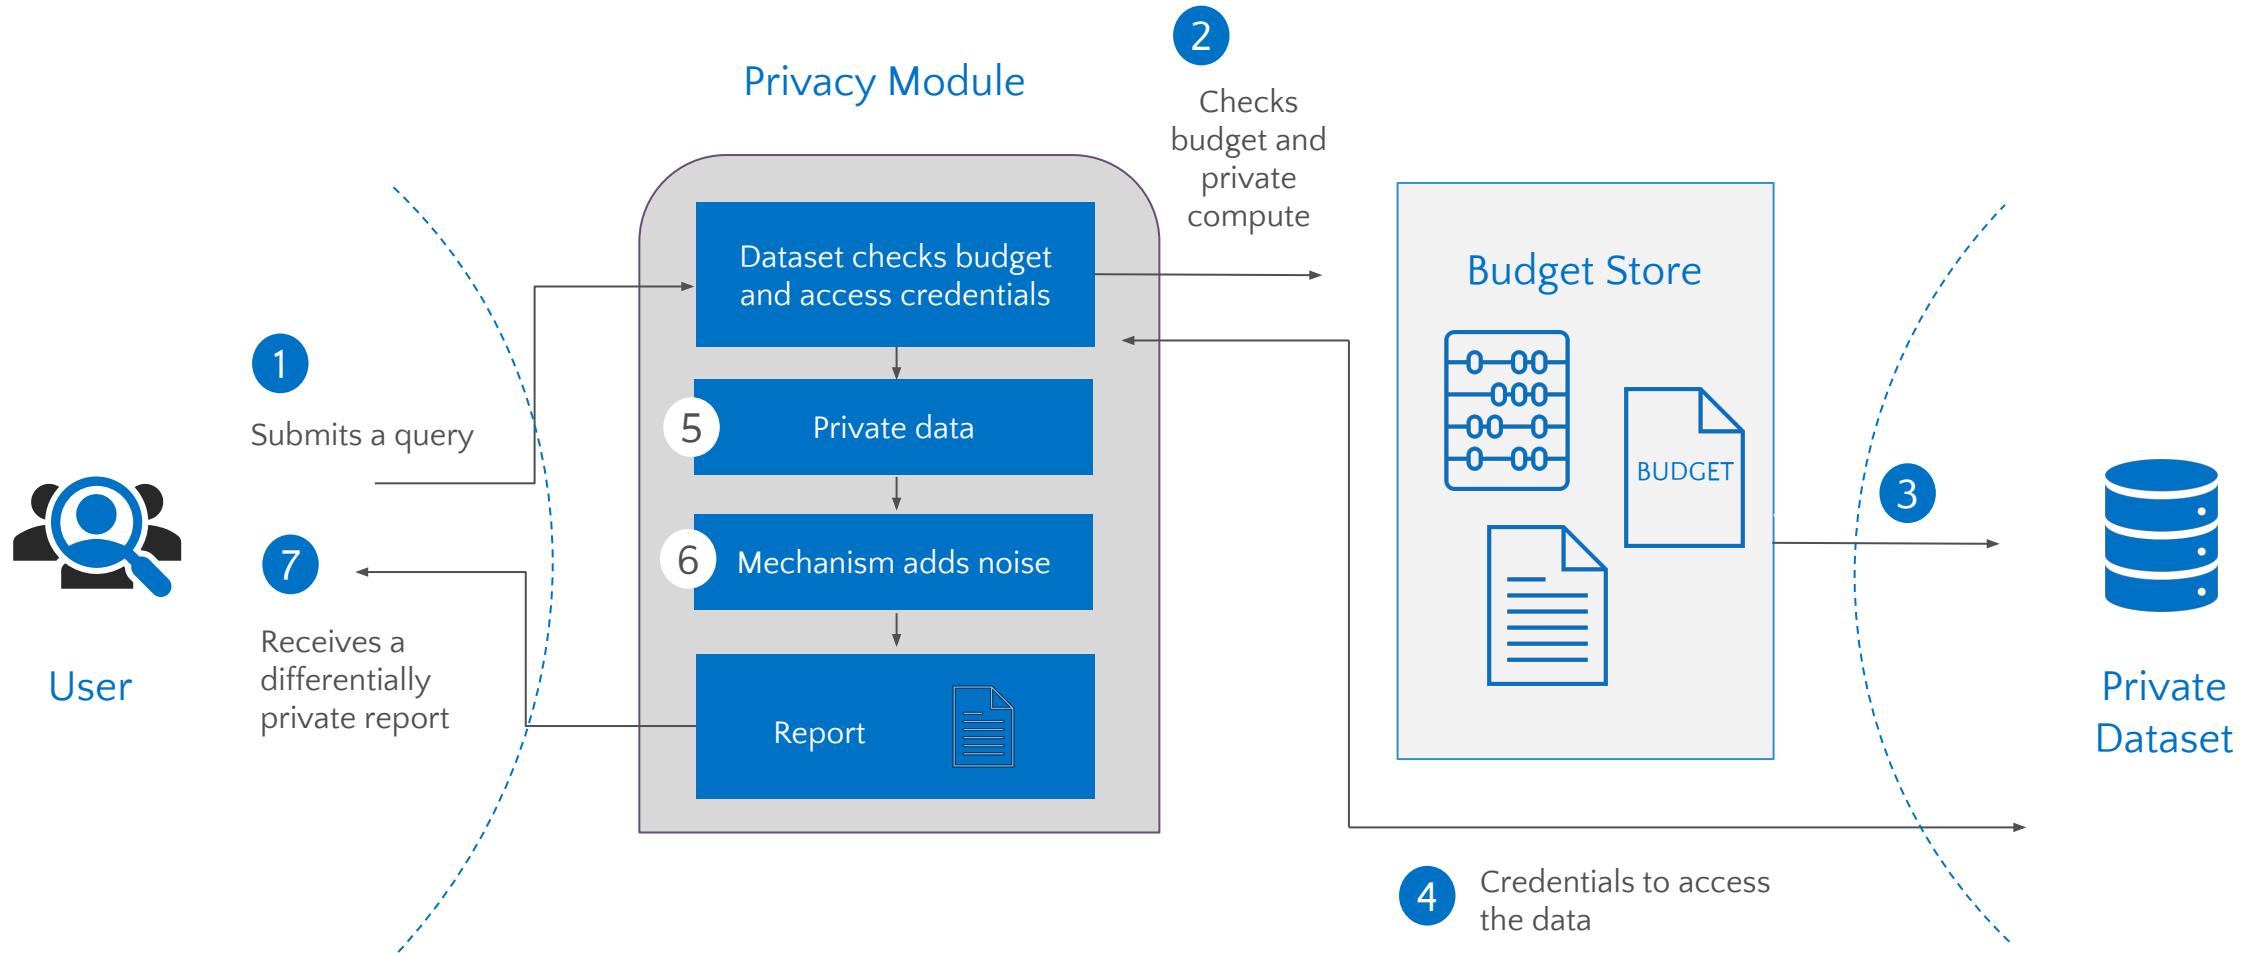
- Future releases will focus on hardened scenarios where the researcher or analyst is untrusted and does not have access to the data directly

<https://opendifferentialprivacy.github.io>

<https://github.com/opendifferentialprivacy>



# WhiteNoise



<https://opendifferentialprivacy.github.io>  
<https://github.com/opendifferentialprivacy>

 **WhiteNoise**

# WhiteNoise

## WhiteNoise Core

- A pluggable open source library of differentially private algorithms and mechanisms with implementations based on mature differential privacy research
- APIs for defining an analysis and a validator for evaluating these analyses and composing the total privacy loss on a dataset

WhiteNoise handles managing the budget for each query and adding the appropriate amount of noise-based budget

## WhiteNoise System

System components to make it easier to interface with your data systems, including a SQL query dialect and connections to many common data sources, such as PostgreSQL, SQL Server, Spark, Presto, Pandas, and CSV files

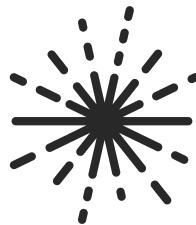
## WhiteNoise Samples

Example code and notebooks to demonstrate the use of the WhiteNoise platform, teach the properties of differential privacy, and highlight some of the nuances of the WhiteNoise implementation



<https://opendifferentialprivacy.github.io>

<https://github.com/opendifferentialprivacy>



# OpenDP

Building an open source suite of tools for deploying differential privacy

A community effort in partnership with Microsoft to build a trustworthy and open-source suite of differential privacy tools that can be easily adopted by custodians of sensitive data to make it available for research and exploration in the public interest

Incubated by Harvard University's [Privacy Tools](#) and [Privacy Insights](#) projects (at [SEAS](#) and [IQSS](#))

With stakeholders and contributors from across academia, industry, and government

Plan to design, implement, and govern an "OpenDP Commons" that includes a library of differentially private algorithms and other general-purpose tools for use in end-to-end differential privacy systems

<https://projects.iq.harvard.edu/opendp>



HARVARD  
UNIVERSITY

# To go beyond...

Differential Privacy: series of lecture from Cynthia Dwork (Microsoft Research)

[Lecture 1](#), [Lecture 2](#), [Lecture 3](#), and [Lecture 4](#)

**WhiteNoise project:**

Introducing WhiteNoise: the new differential privacy platform from Microsoft and Harvard's OpenDP

<https://cloudblogs.microsoft.com/opensource/2020/05/19/new-differential-privacy-platform-microsoft-harvard-opendp/>

Preserve data privacy by using differential privacy and the WhiteNoise package

<https://docs.microsoft.com/azure/machine-learning/concept-differential-privacy>

Use differential privacy in Azure Machine Learning

<https://docs.microsoft.com/azure/machine-learning/how-to-differential-privacy>

WhiteNoise: A Platform for Differential Privacy

<Https://aka.ms/WhiteNoiseWhitePaper>